

CS 576 – Systems Security

Smashing the Stack

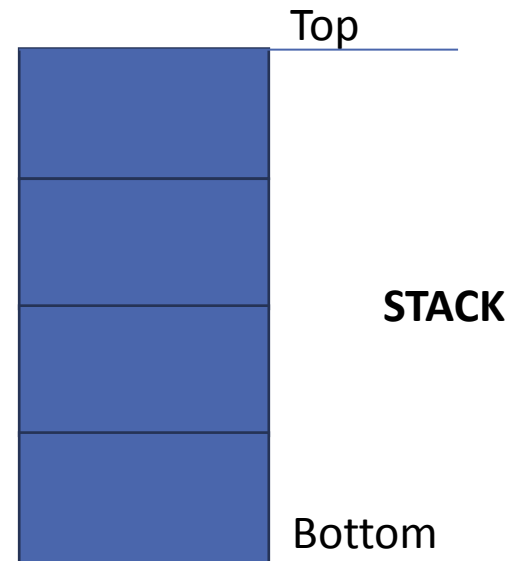
Georgios (George) Portokalidis




How Do Function Calls Work

Stack Data Structure

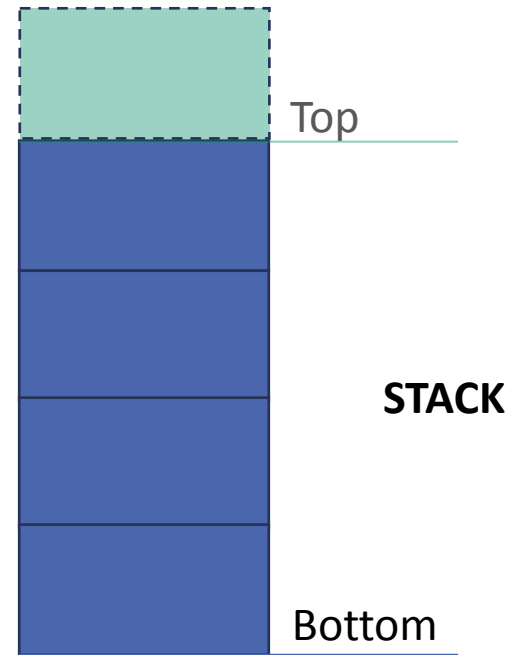
- Stack plays a crucial role in supporting functions
 - Follows last-in first-out semantic




Stack Operation Push




Stack Operation Push




Stack Operation Push




Stack Operation Pop



Stack Operation Pop




Stack Operation Pop



The Stack Pointer (SP)


The stack pointer points to the first element in the stack (the top).



The Stack Pointer (SP)

The stack pointer points to the first element in the stack (the top).

Usually the RSP/ESP register is used to store the SP.




The Stack Pointer (SP)

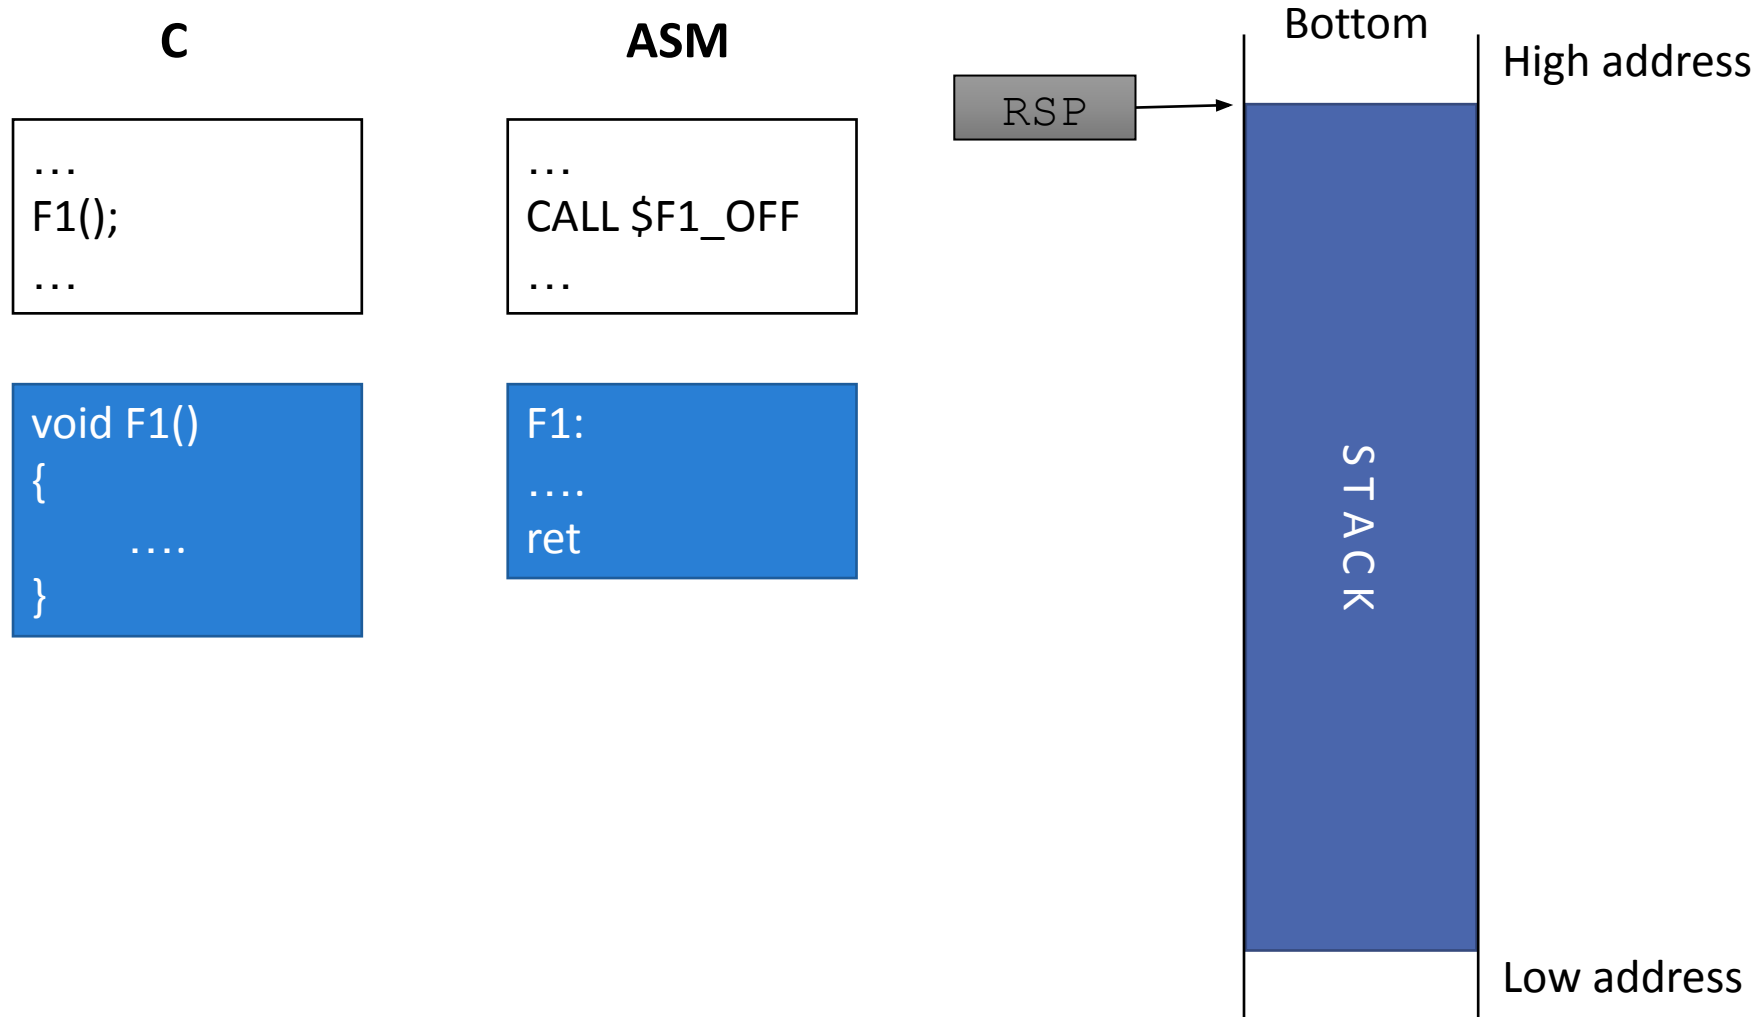
The stack pointer points to the first element in the stack (the top).

Usually the RSP/ESP register is used to store the SP.

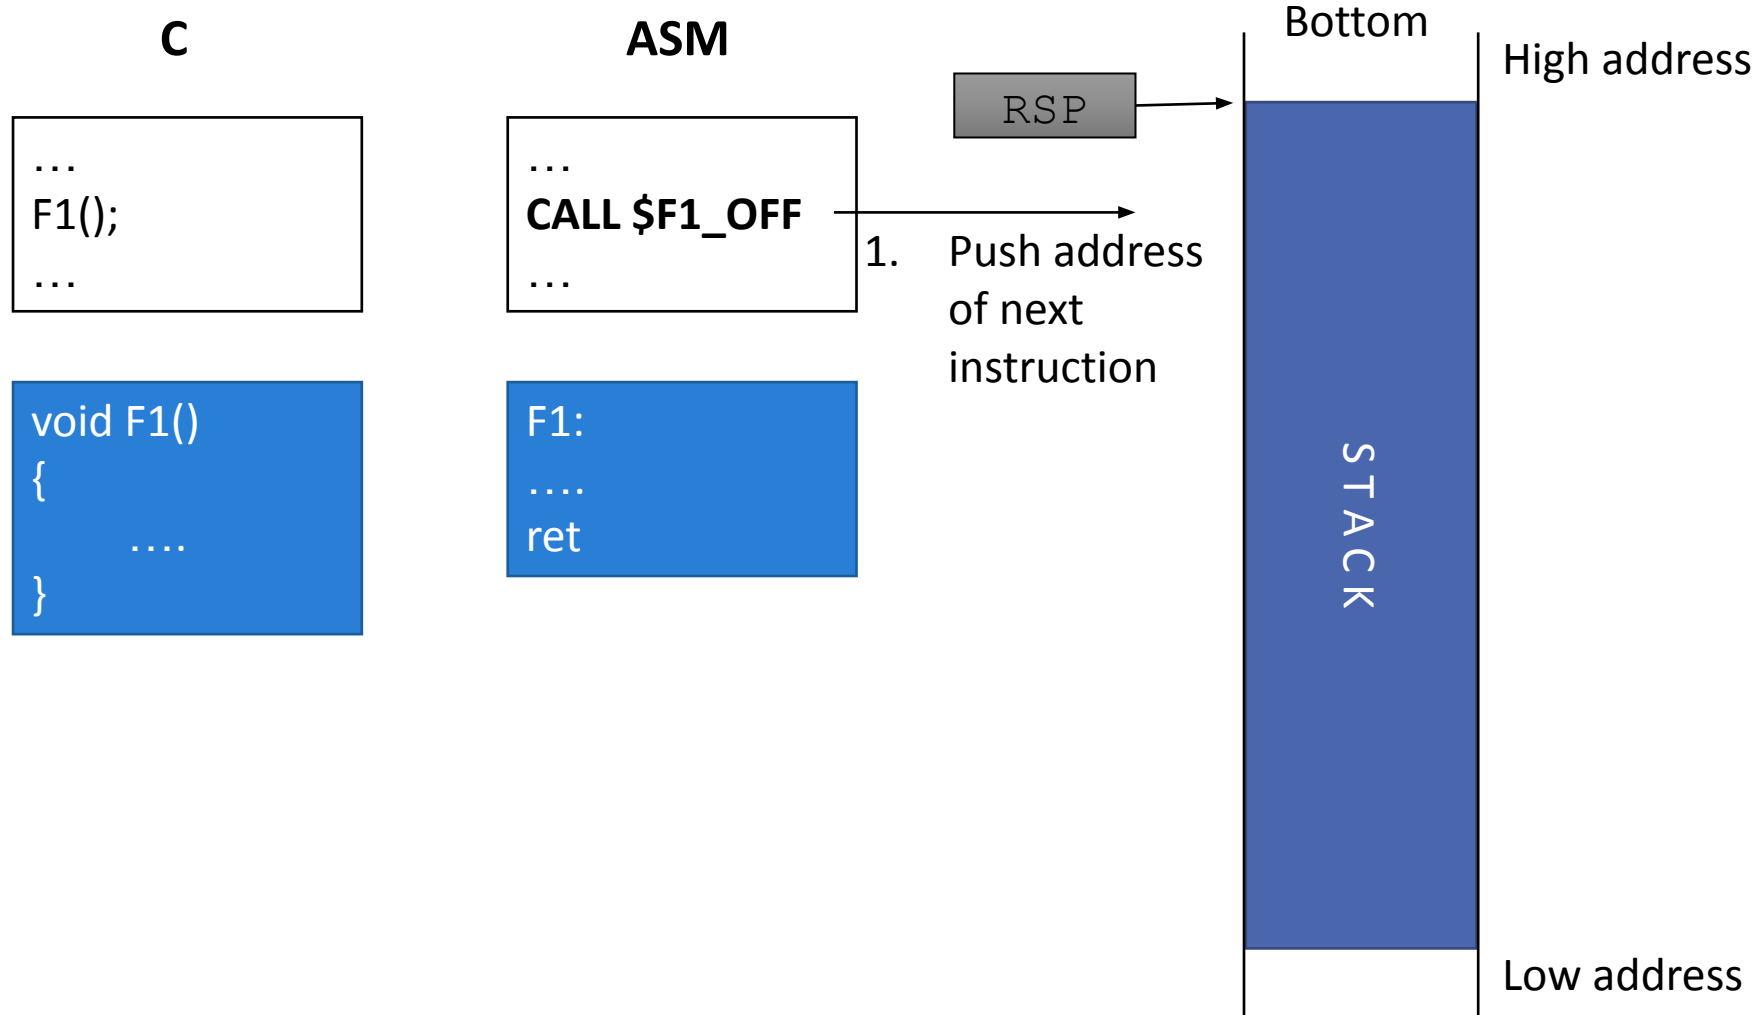
The stack grows towards lower addresses



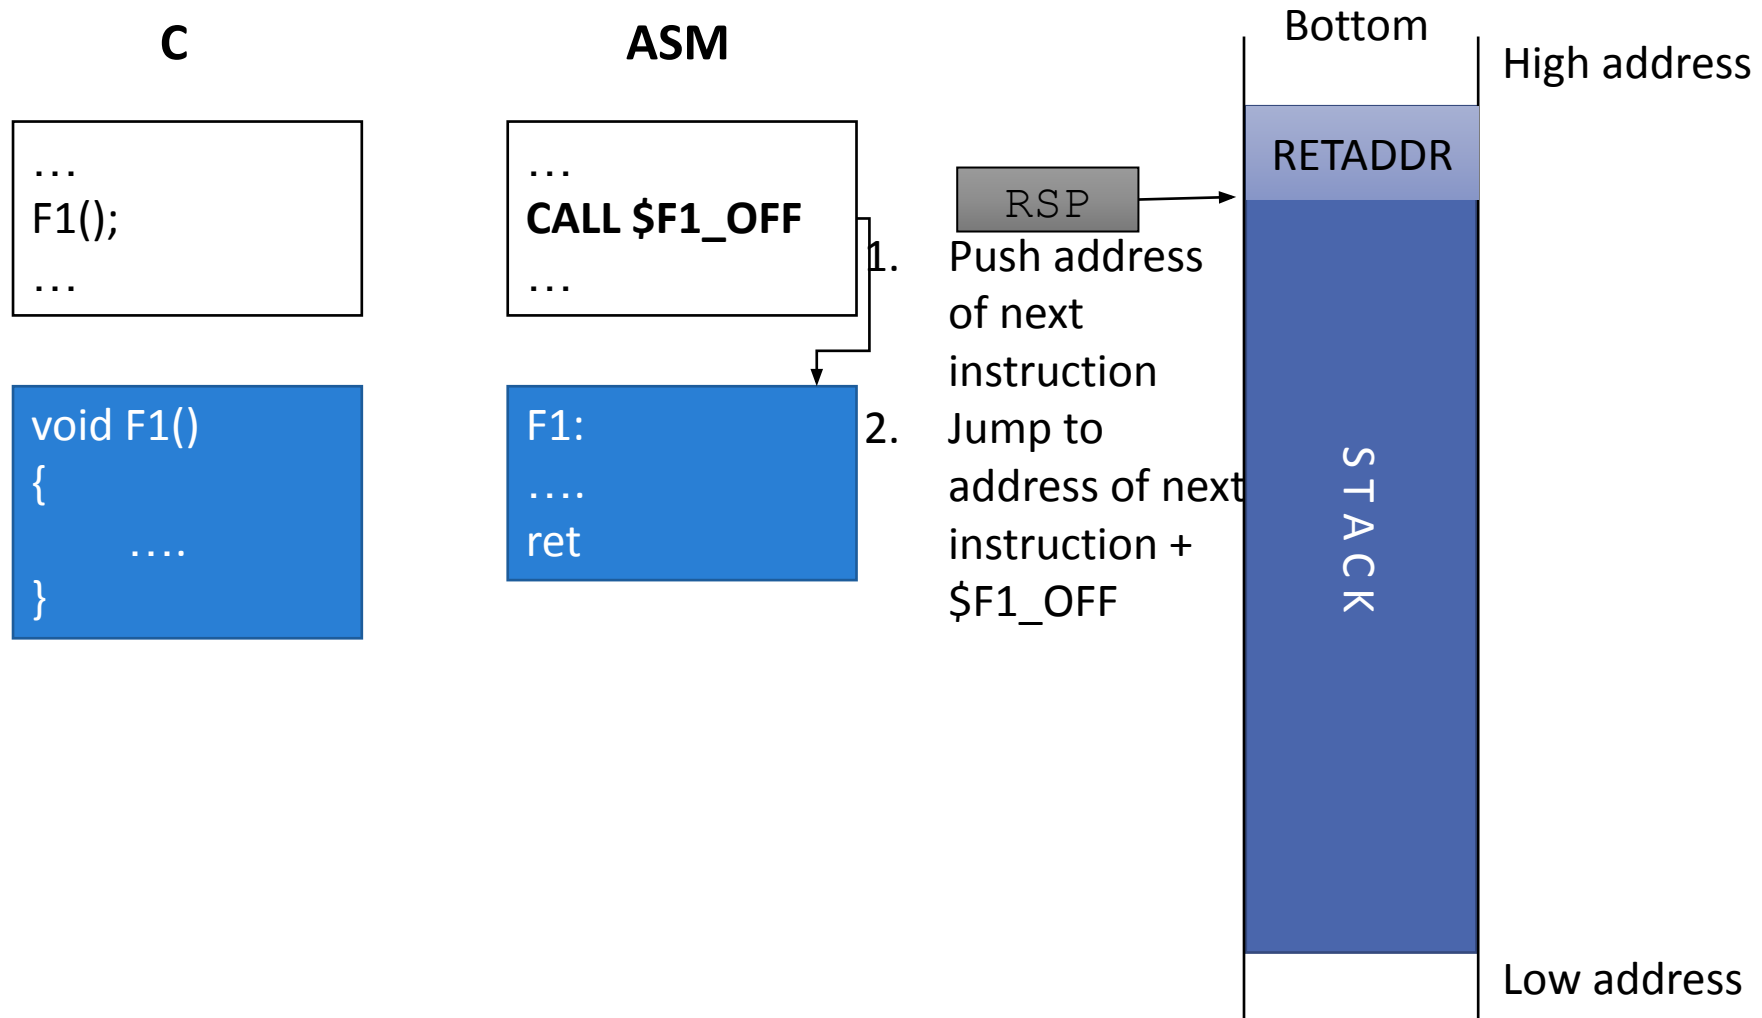
Simple Function Call



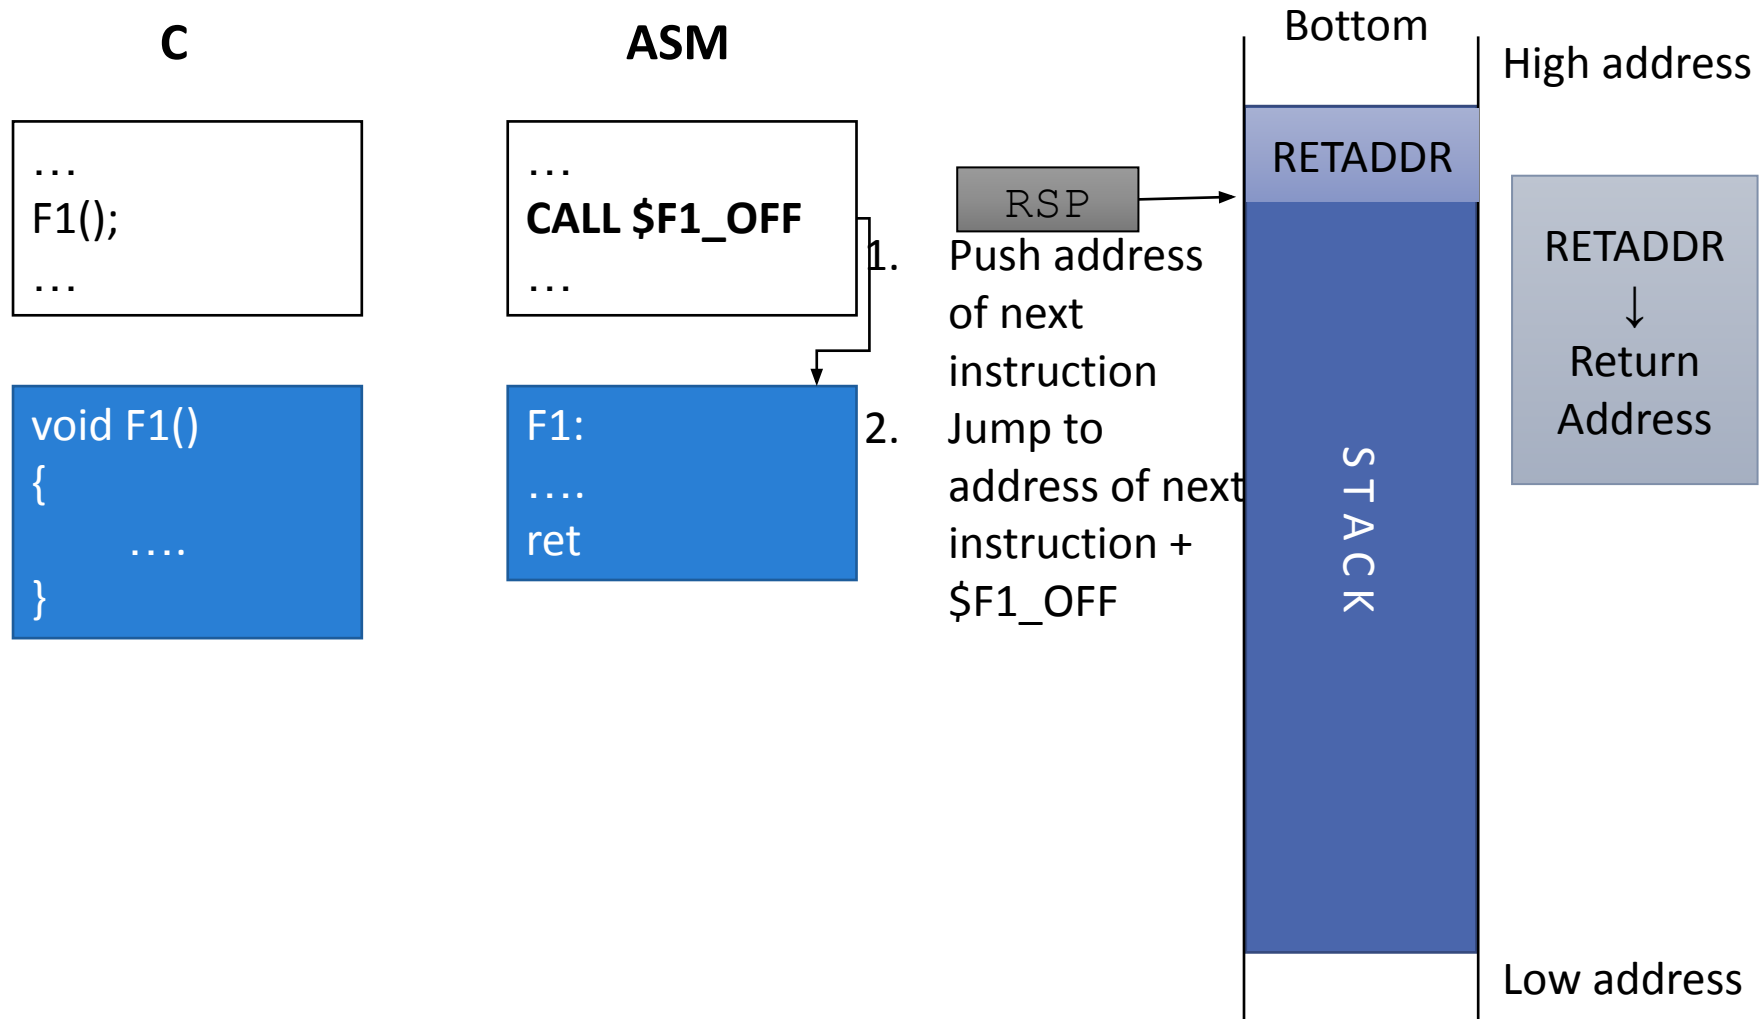
Simple Function Call



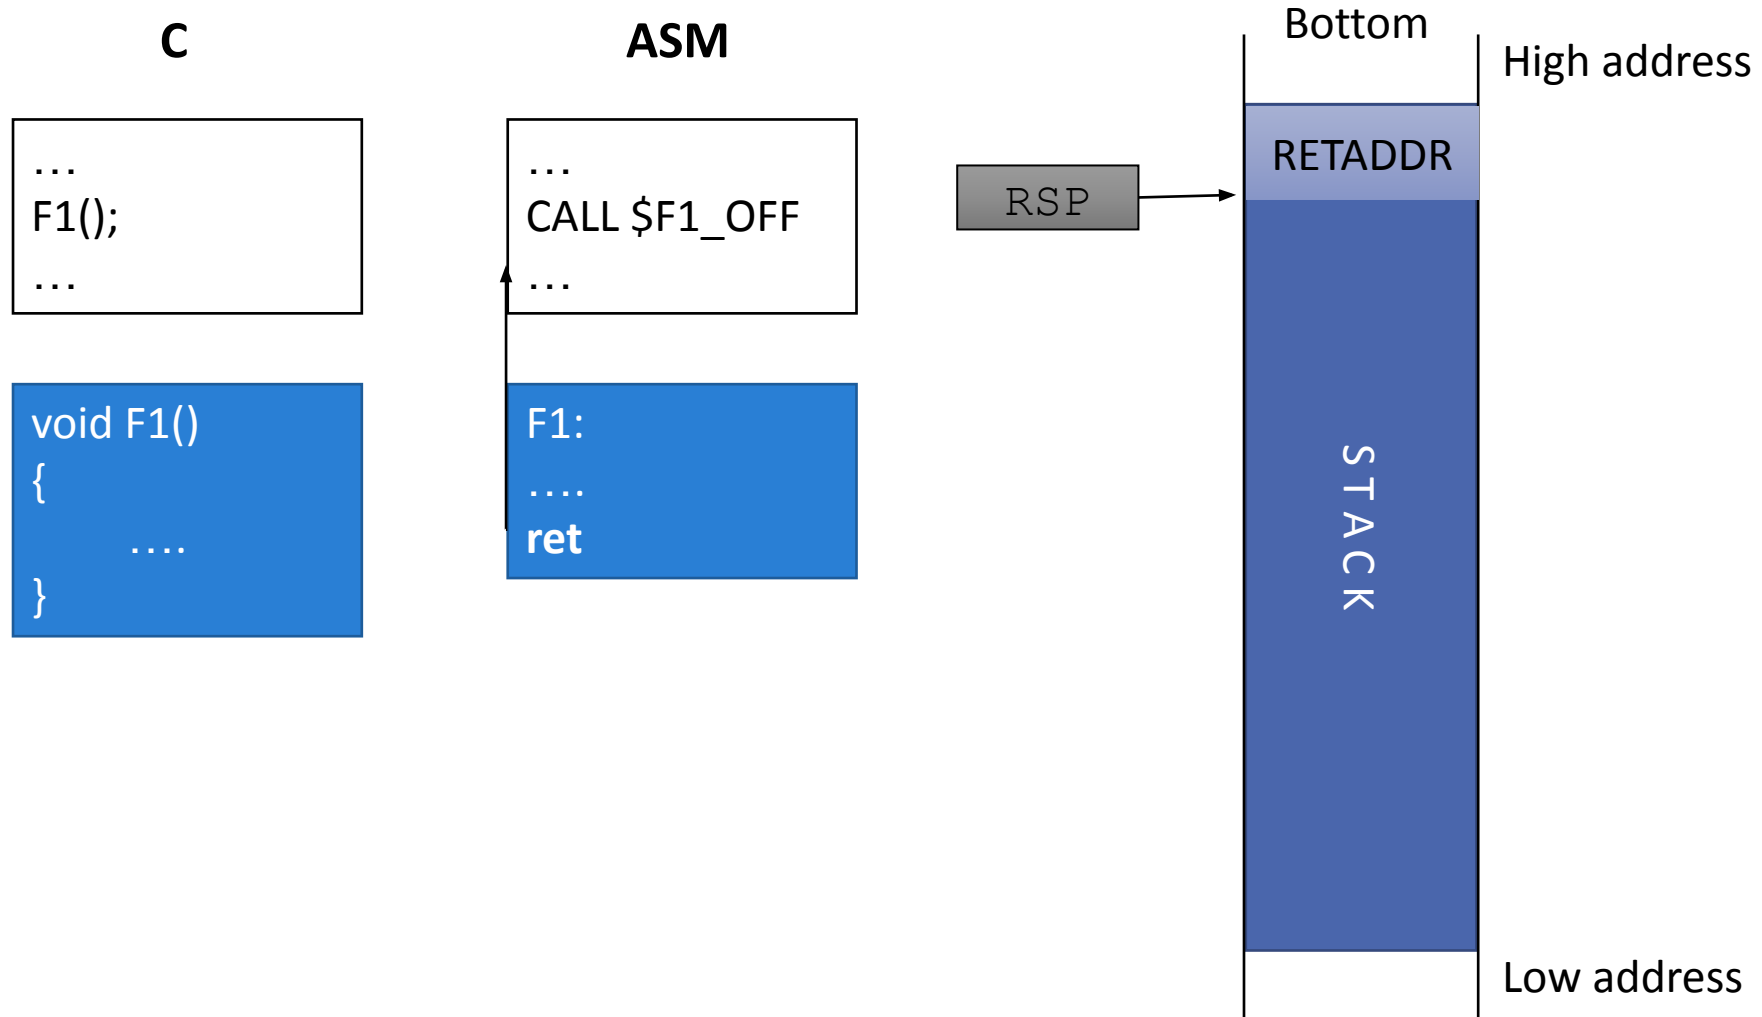
Simple Function Call



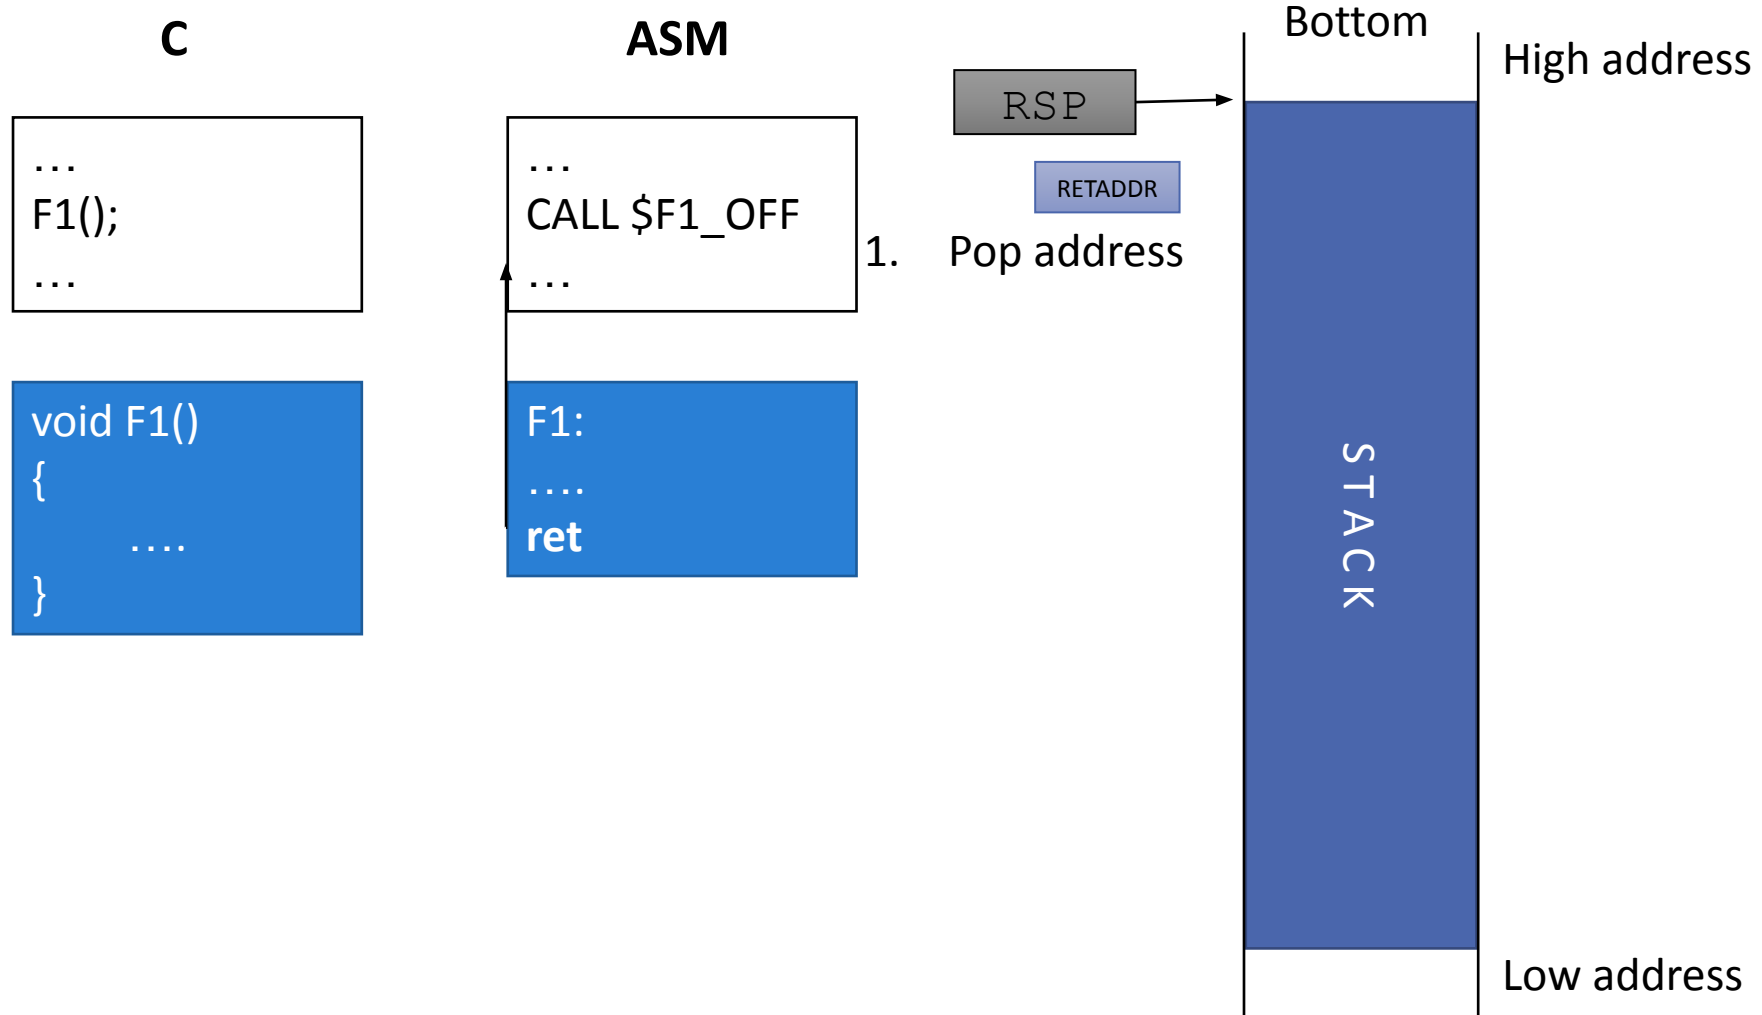
Simple Function Call



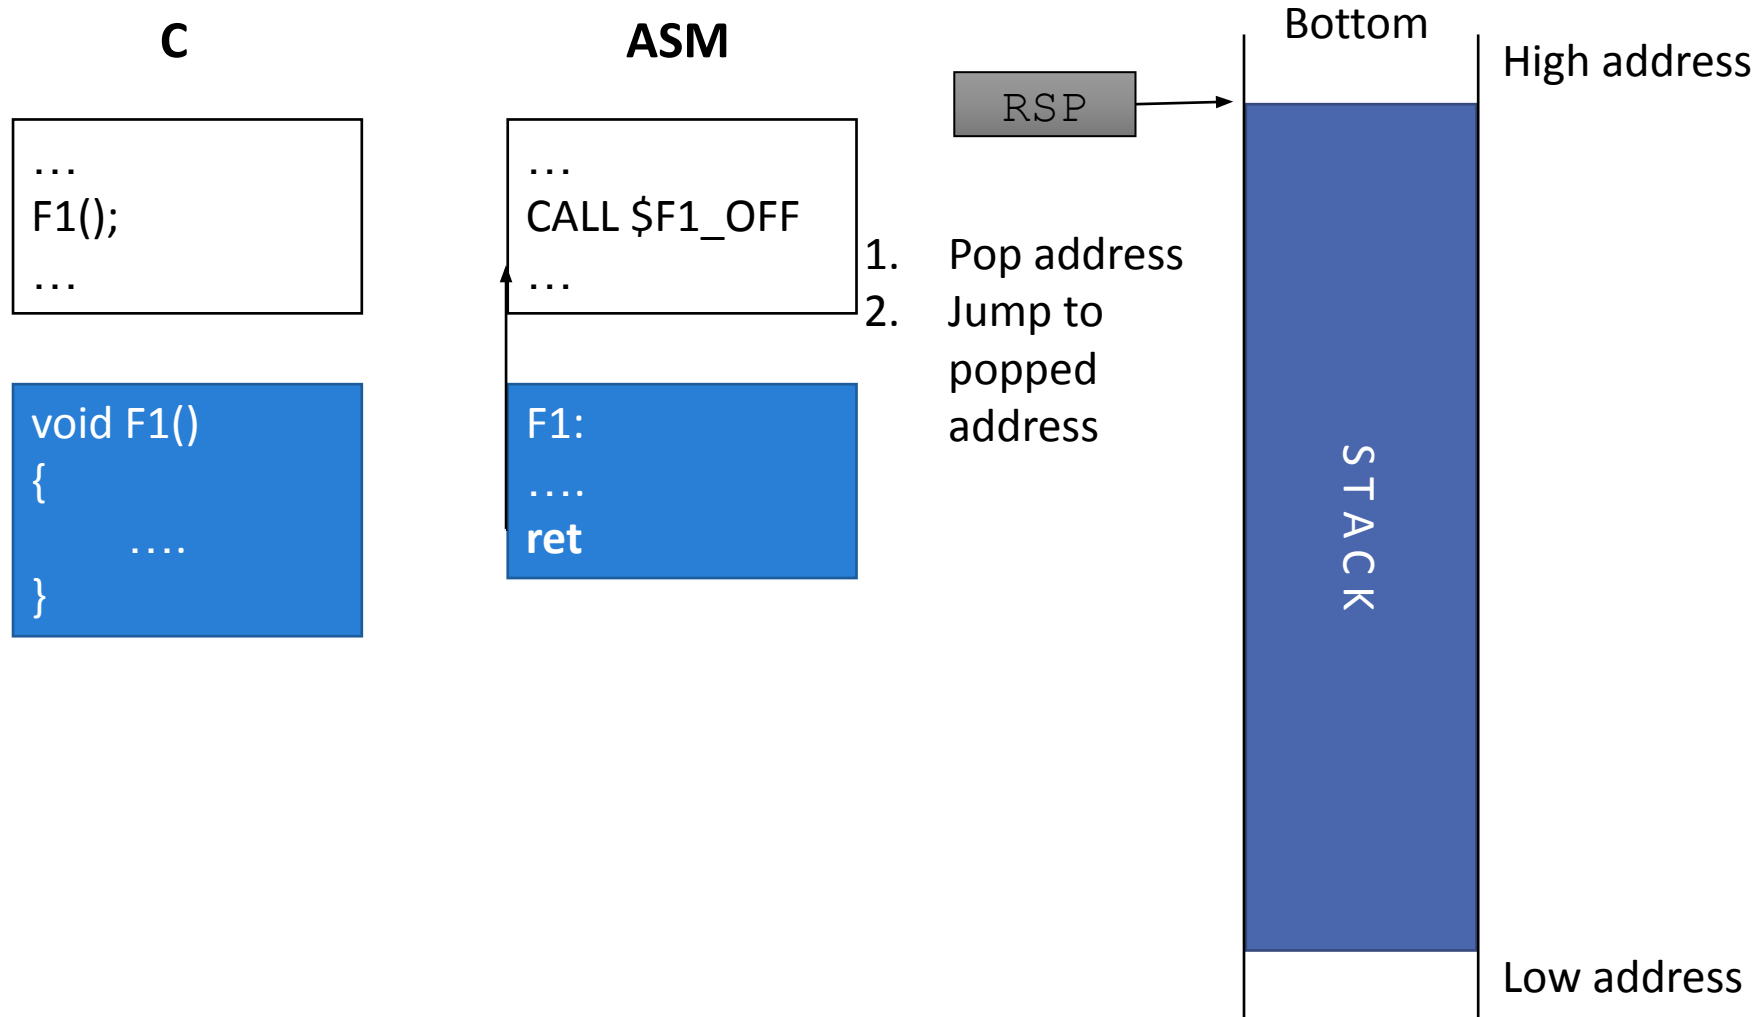
Simple Function Return



Simple Function Return



Simple Function Return



Function Calls and Returns

Calling a function (the callee)

- CALL instruction
 - Pushes `next_ins_addr` on stack and transfers control to address described by operand
- Most common syntax: CALL OFFSET
 - Target is `next_ins_addr + OFFSET`

Returning to caller


- RET instruction
 - Pop return address from stack and transfers control to it

```
CALL tgt □ push next_ins_addr; jmp tgt  
RET □ pop retaddr; jmp retaddr
```

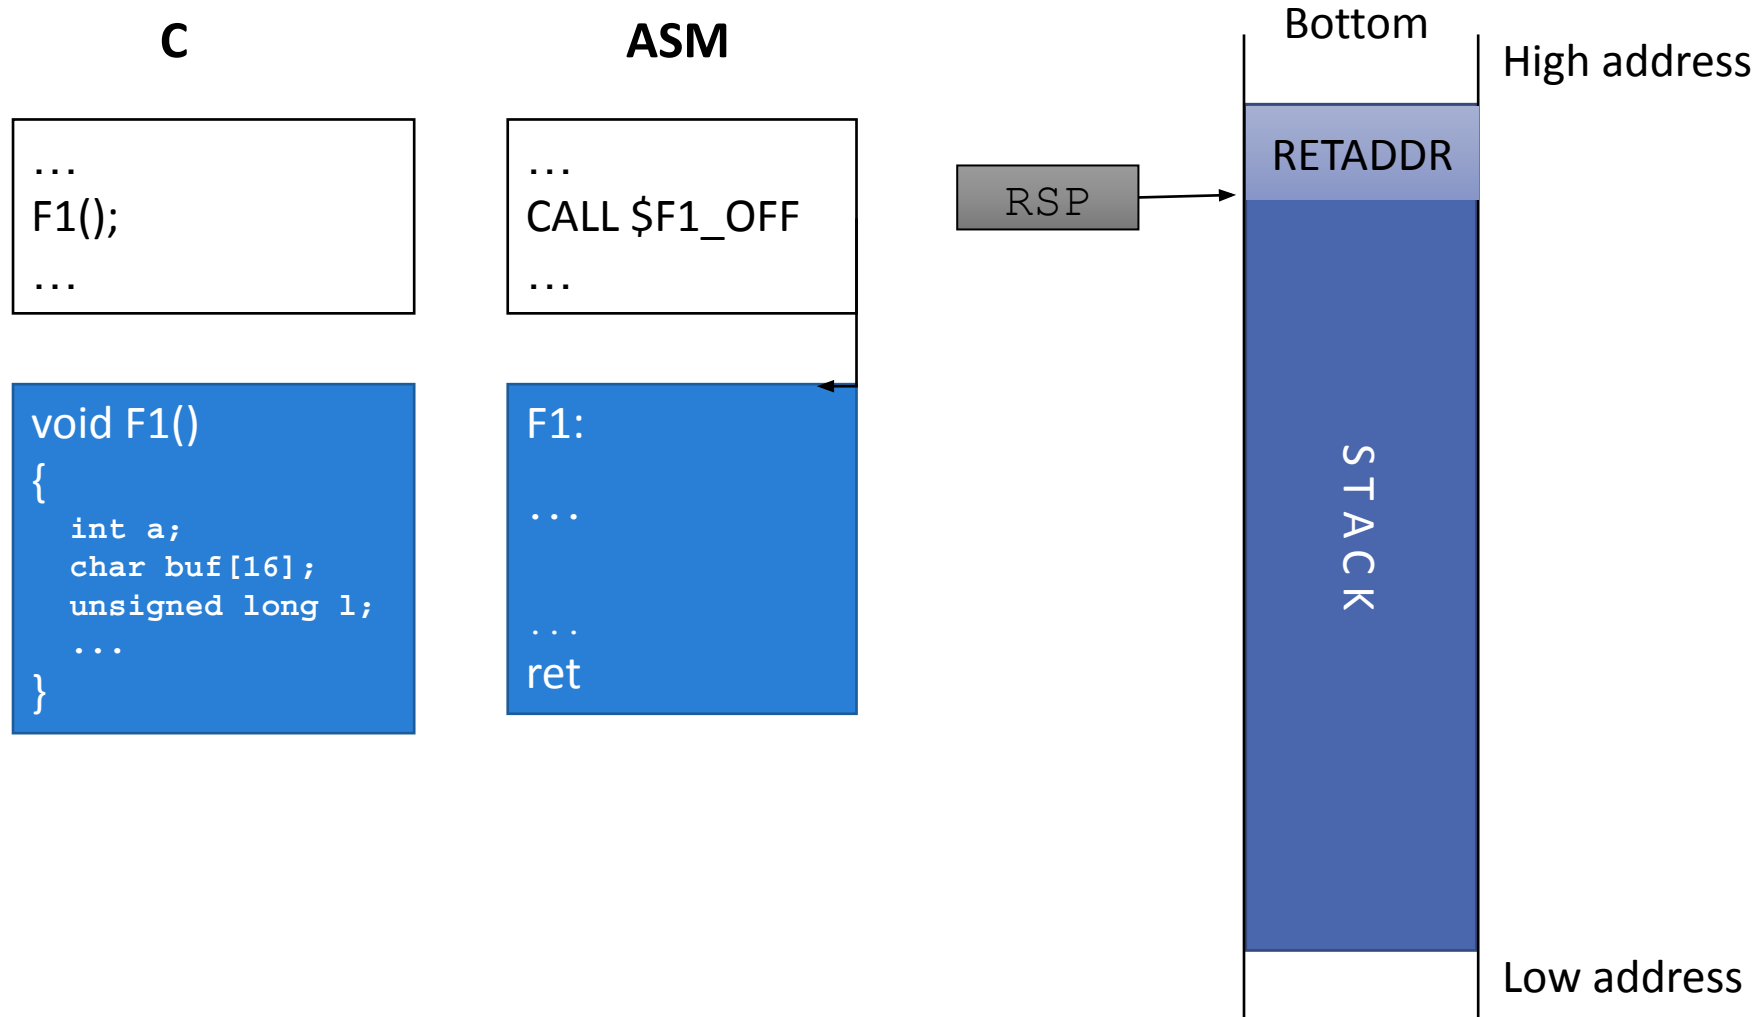
```
call and ret implicitly use the SP
```

The Stack Is Used...

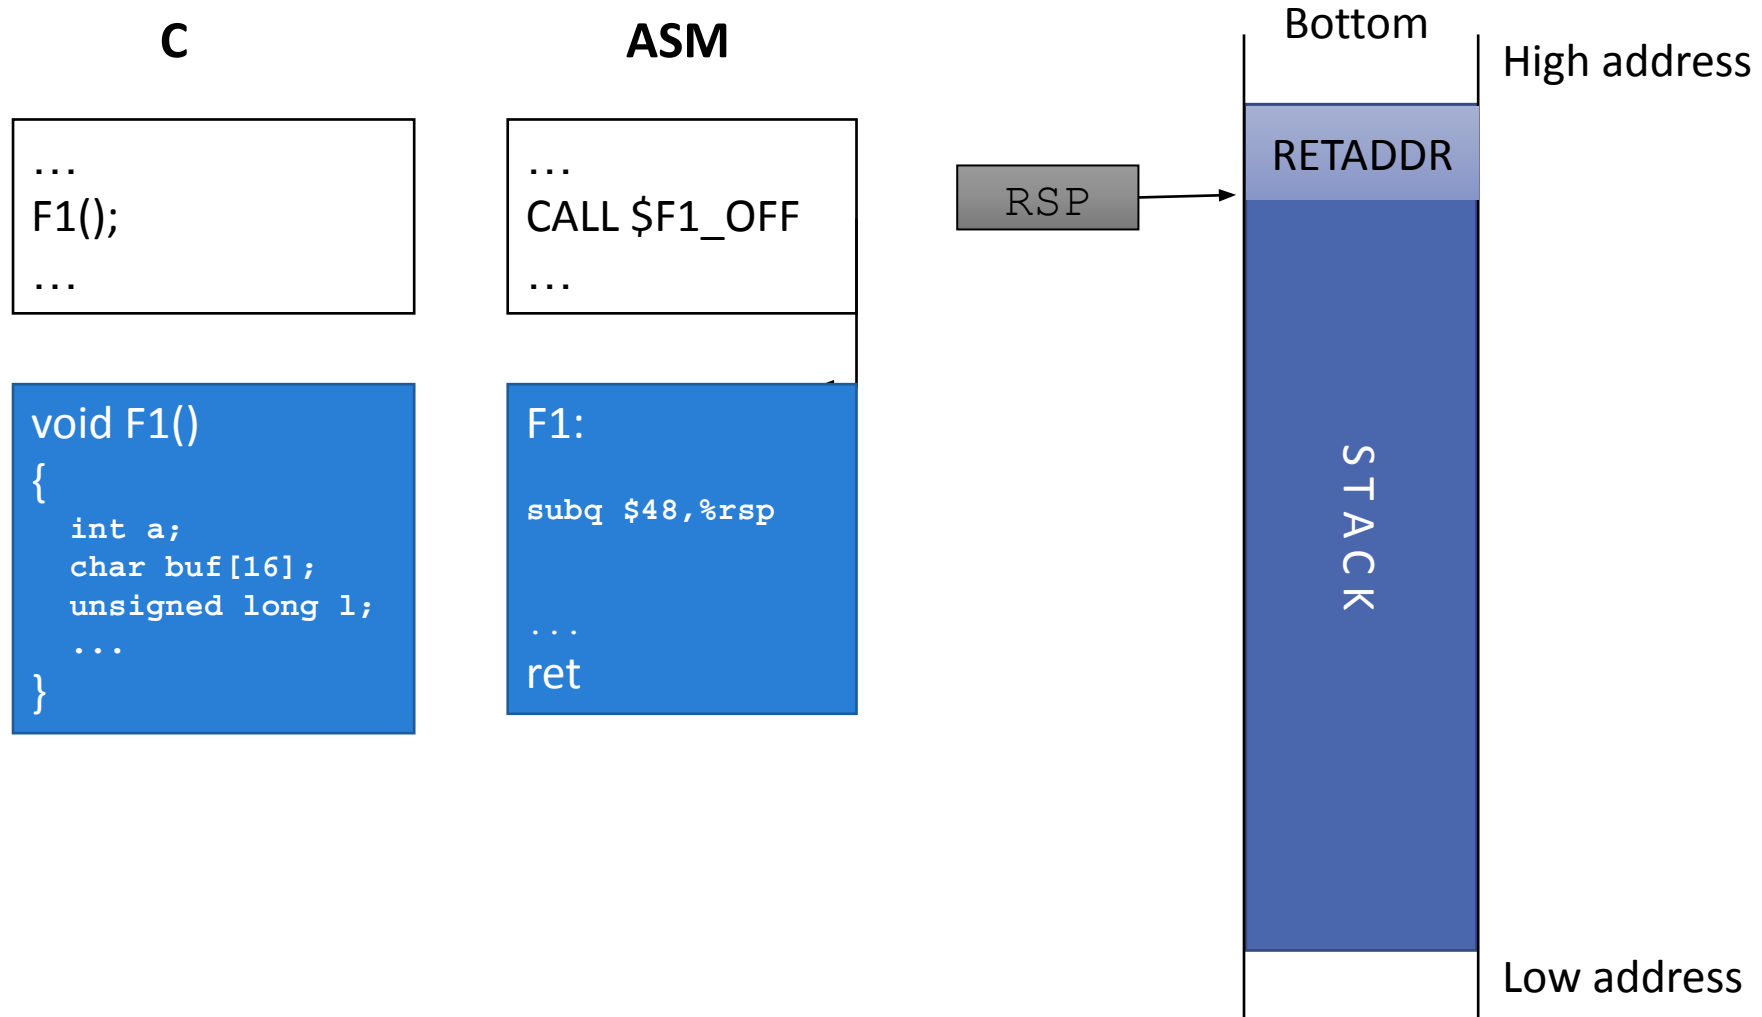
- ...to store the return address of caller functions
 - Code pointers!
- ...to store local variables
 - Aka stack variables



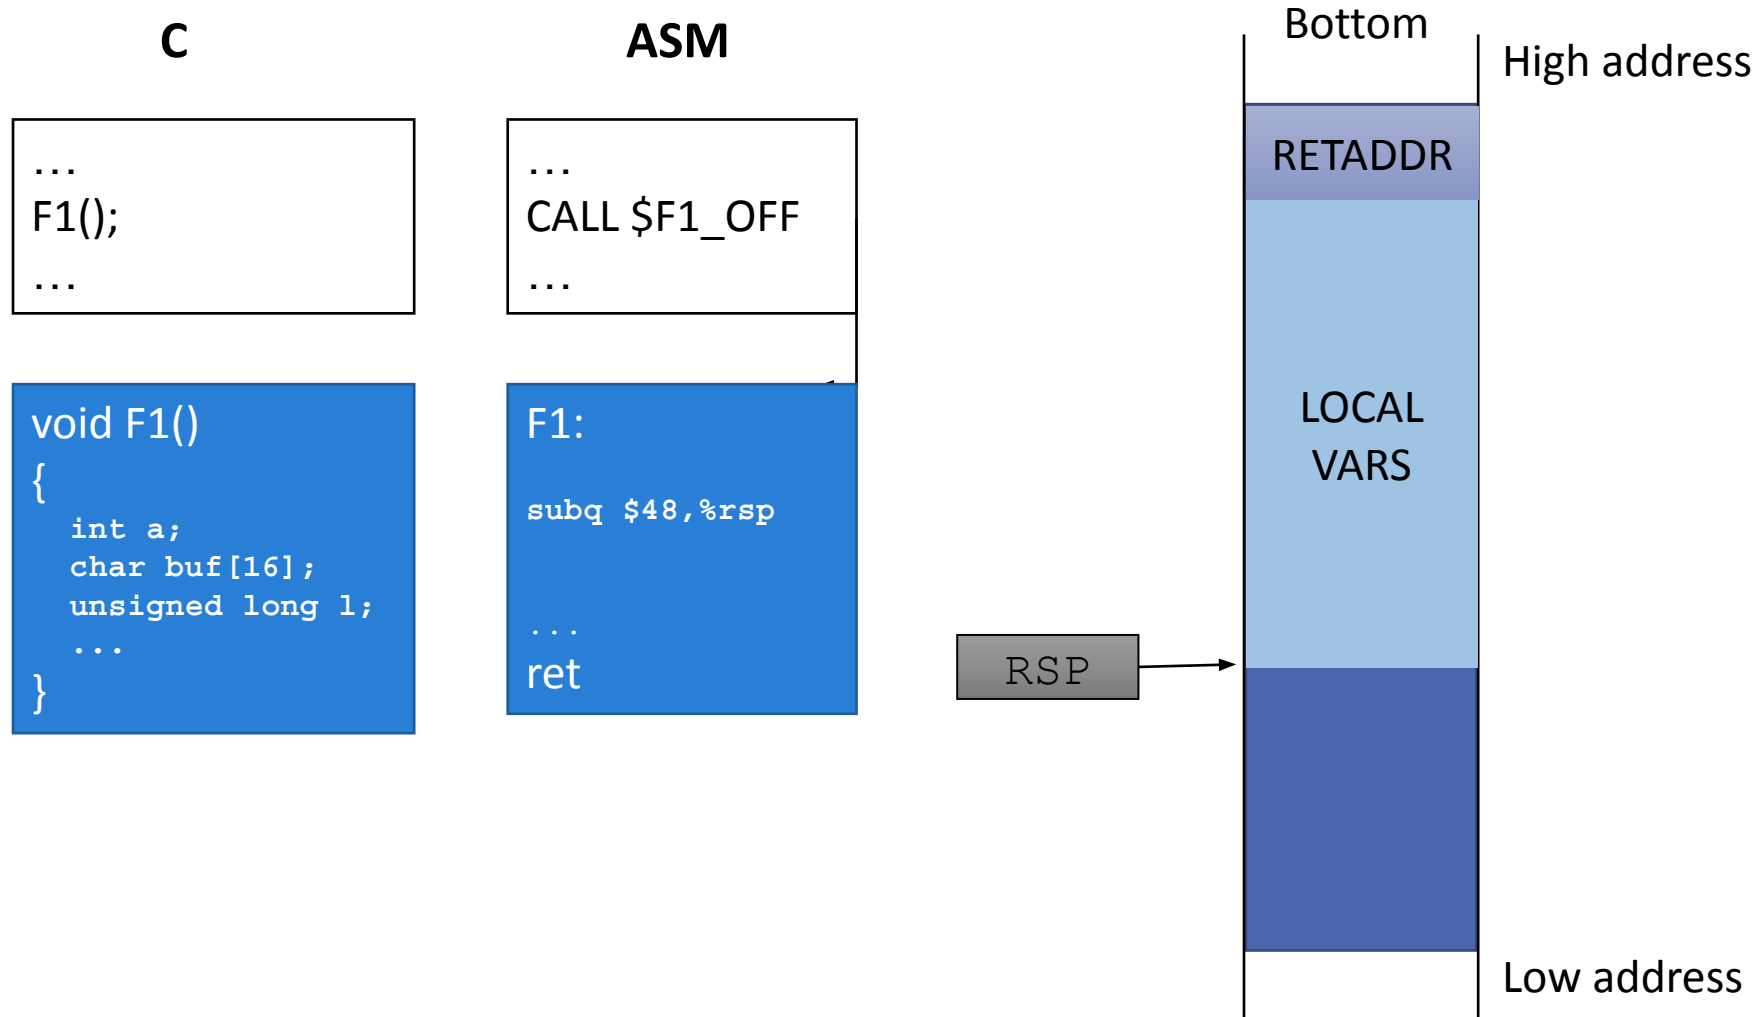
Local Variables



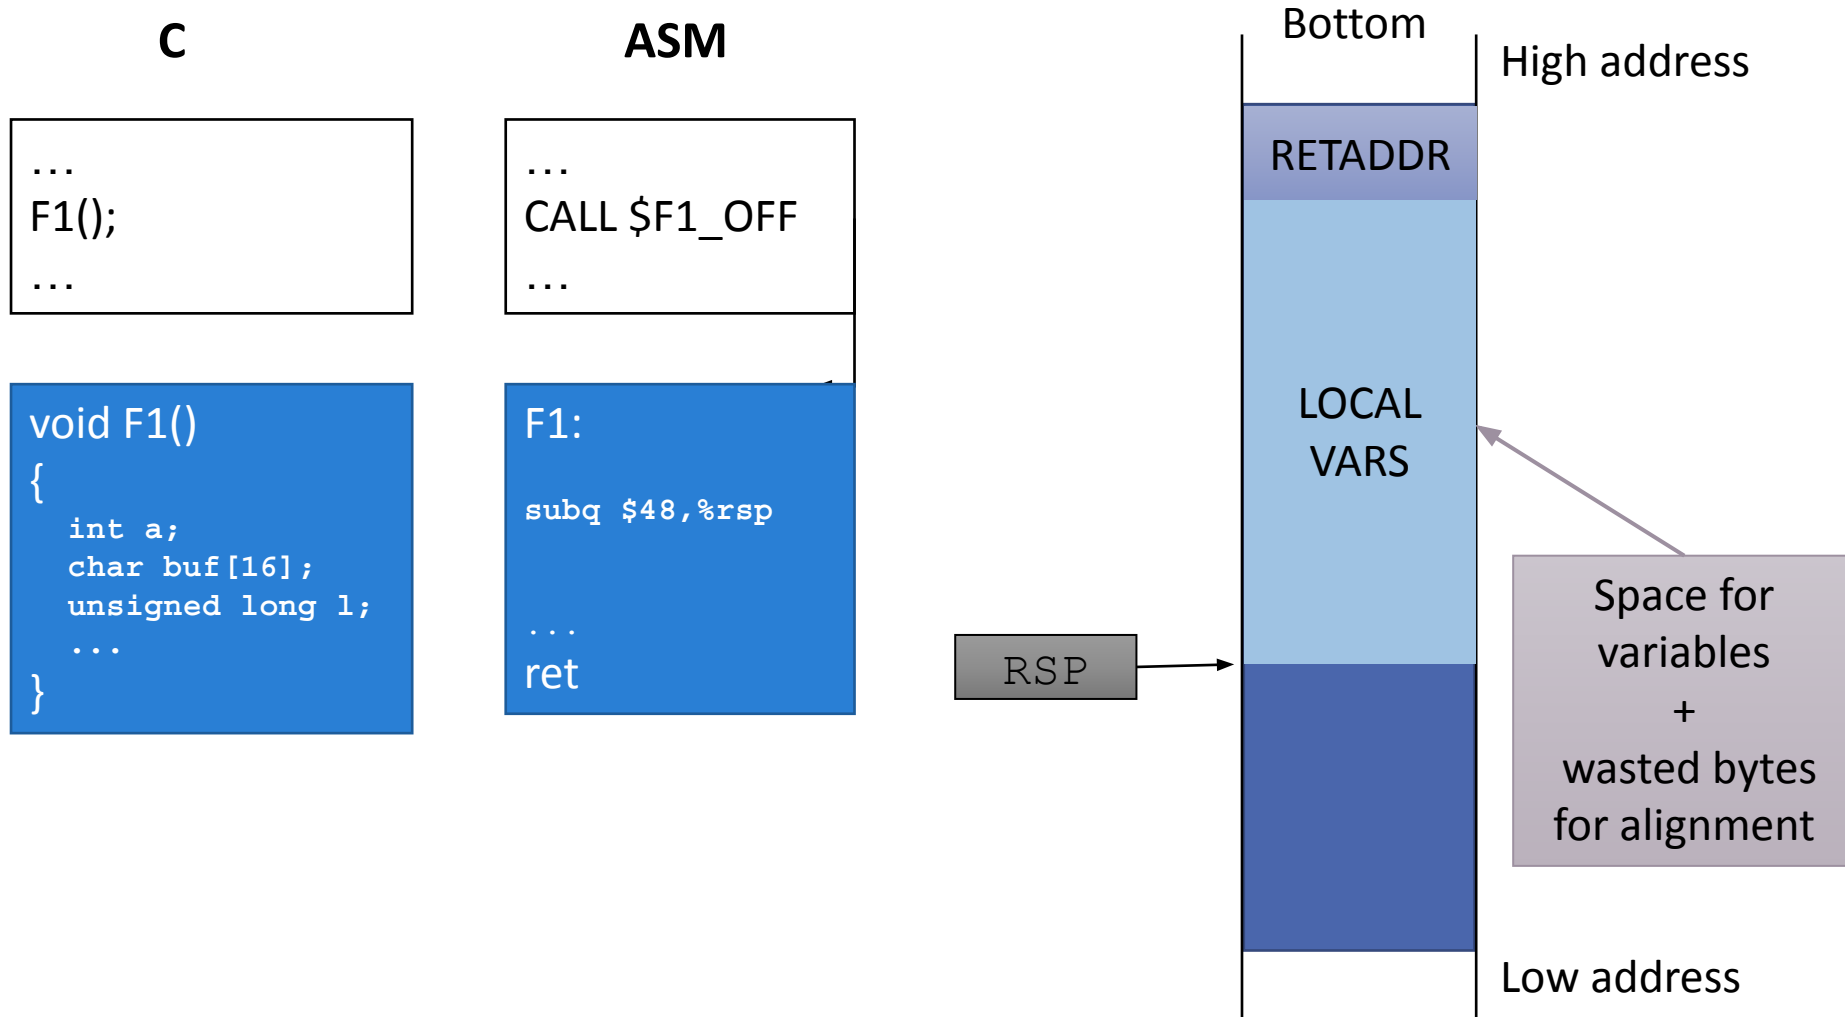
Local Variables



Local Variables




Local Variables




Alignment

- CPUs like aligned data
 - Better performance
- Compilers try to align data



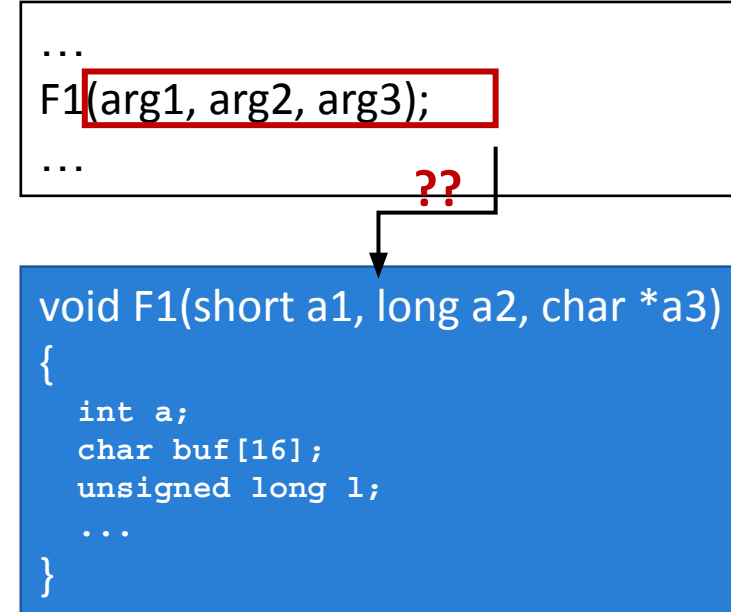
The Stack Is Used...

- ...to store the return address of caller functions
 - Code pointers!
- ...to store local variables
 - Aka stack variables
- ...to pass function arguments



Calling Conventions

- Defines the standard for passing arguments
- Caller and callee need to agree on the standard
- Enforced by compiler
- Important when using 3rd party libraries
 - Hence, also referred to as the Application Binary Interface (ABI)
- Different styles ↔ different advantages



Popular conventions:


- cdecl (32-bit)
- System V AMD64 ABI

cdecl

- Arguments are passed on the stack
 - Pushed right to left

```
...  
F1(0xff, UINT_MAX, argv[0]);  
...
```

```
...  
pushl    (%eax)  
pushl    $-1  
pushl    $255  
call     F1  
...
```




cdecl

- Arguments are passed on the stack
 - Pushed right to left


```
...  
F1(0xff, UINT_MAX, argv[0]);  
...
```

```
...  
pushl    (%eax)  
pushl    $-1  
pushl    $255  
call     F1  
...
```



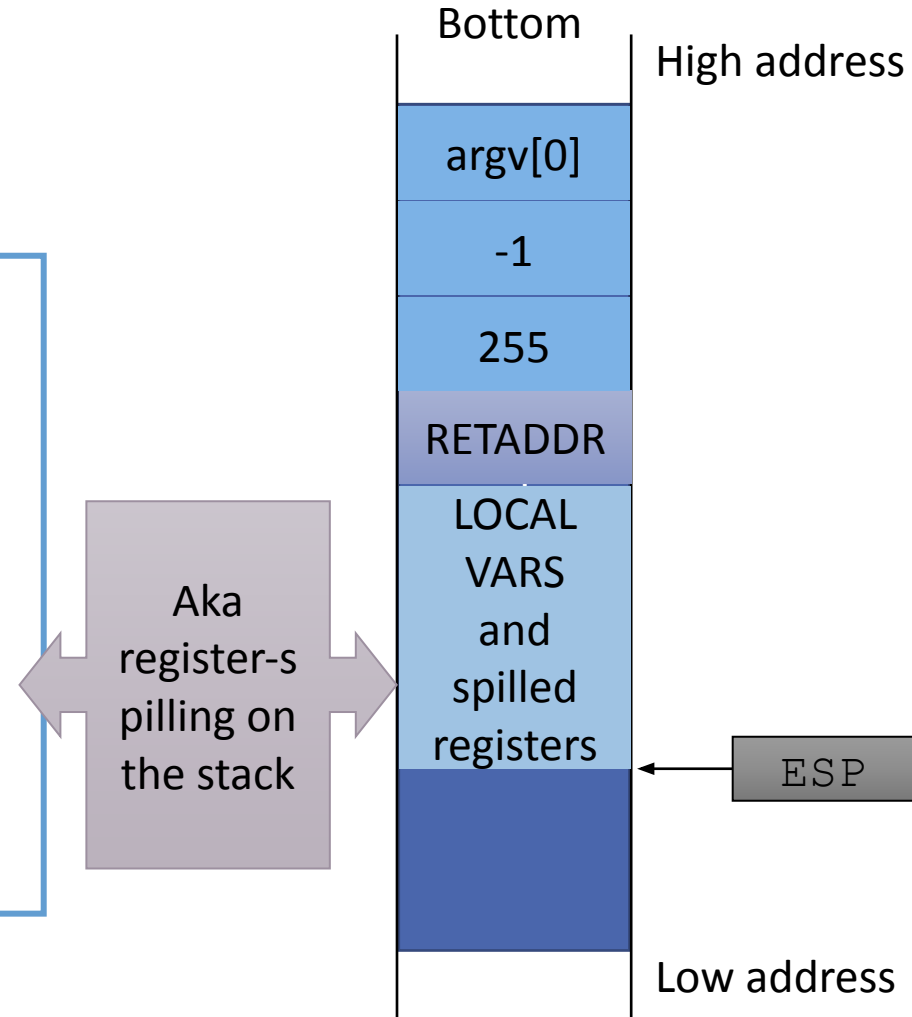
cdecl

- Arguments are passed on the stack
 - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
 - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
 - callee must ensure they have same value on return



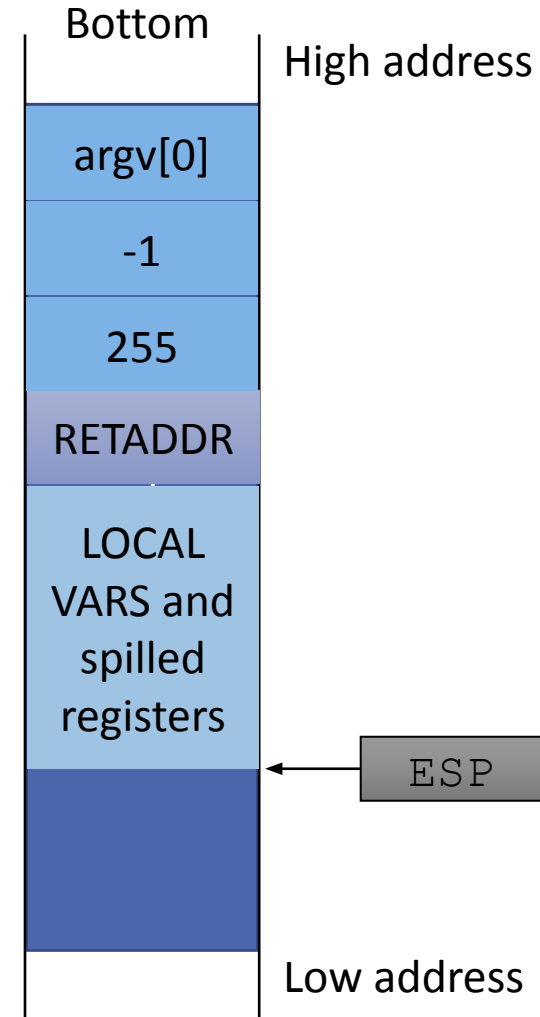
cdecl

- Arguments are passed on the stack
 - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
 - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
 - callee must ensure they have same value on return



cdecl

- Arguments are passed on the stack
 - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
 - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
 - callee must ensure they have same value on return
- `eax` used for function return value




System V AMD64 ABI

- Arguments are passed using registers
 - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9

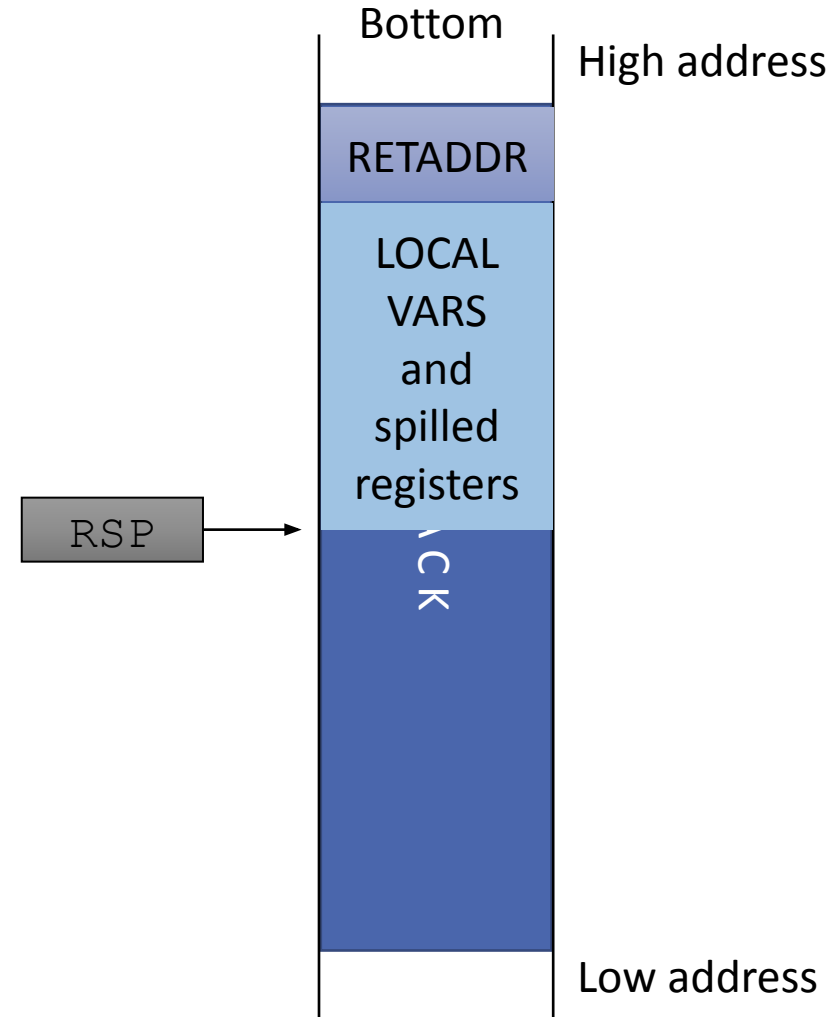
```
...  
F1(0xff, UINT_MAX, argv[0]);  
...
```

```
...  
movq    (%rsi), %rdx  
movl    $4294967295, %esi  
movl    $255, %edi  
call    F1  
...
```




System V AMD64 ABI

- Arguments are passed using registers
 - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9
- RBP, RBX, and R12–R15 are callee saved



System V AMD64 ABI

- Arguments are passed using registers
 - First 6 integer or pointer arguments are passed in registers RDI, RSI, RDX, RCX, R8, and R9
- RBP, RBX, and R12–R15 are callee saved
- RAX used for function return



Popular Conventions Summary

cdecl (mostly 32-bit)

- Arguments are passed on the stack
 - Pushed right to left
- `eax`, `edx`, `ecx` are caller saved
 - callee can overwrite without saving
- `ebx`, `esi`, `edi` are callee saved
 - callee must ensure they have same value on return
- `eax` used for function return value

System V AMD64 ABI

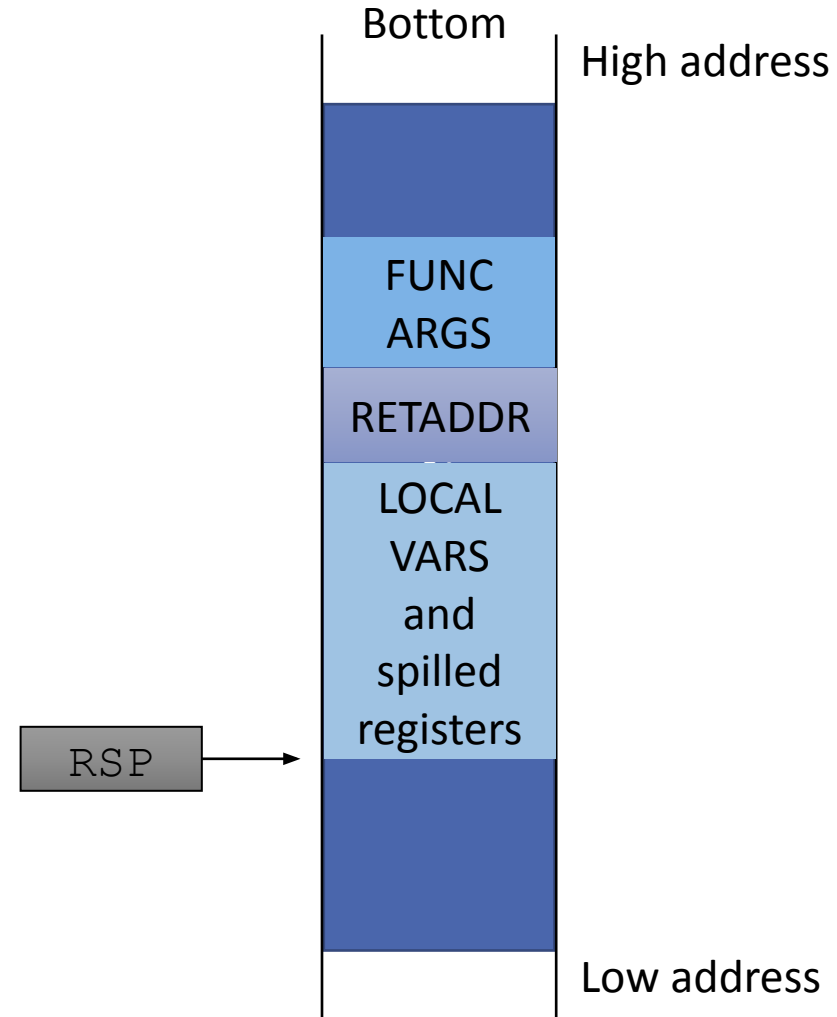
- Arguments are passed using registers
 - First 6 integer or pointer arguments are passed in registers `rdi`, `rsi`, `rdx`, `rcx`, `r8`, and `r9`
- `RBP`, `RBX`, and `R12–R15` are callee saved
- `RAX` used for function return

Conventions include additional information, consult reading material for thorough description

- Example: handling of floating point regs

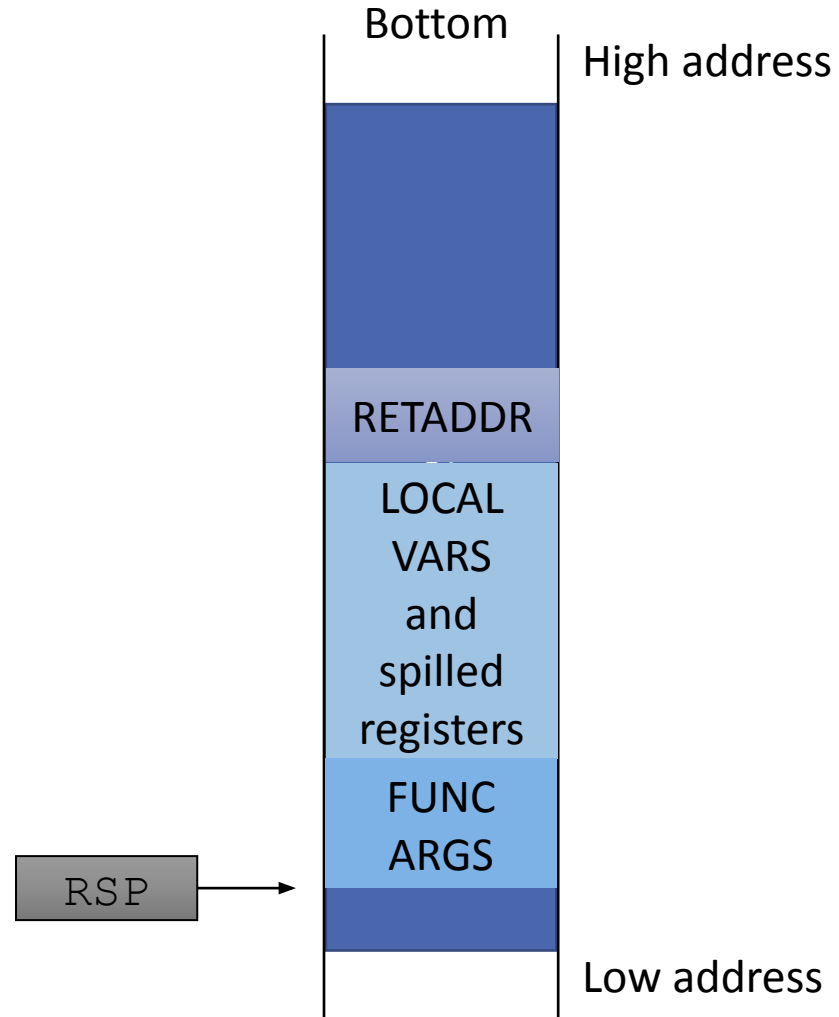
The Stack Is Used...

- ...to store the return address of caller functions
 - Code pointers!
- ...to store local variables
 - Aka stack variables
- ...to pass function arguments
- ...to temporarily store register values
- ...to store the frame pointer



Stack Frame

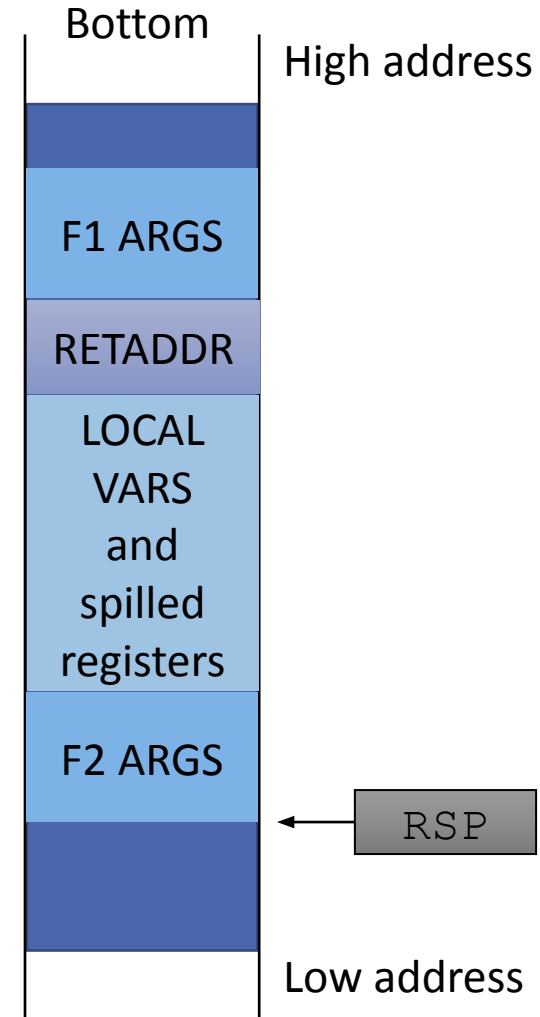
- A stack frame includes all function-local data
 - Local variables
 - Spilled registers
 - Function arguments pushed to the stack to make calls
- More of a logical entity
- Can grow as function executes



Stack Frame Boundaries


- Start below return address
- Stop at stack pointer

```
void F1(short a1, long a2, char *a3)
{
    int a;
    char buf[16];
    unsigned long l;
    ...
    long l2 = F2(a);
    ...
}
```



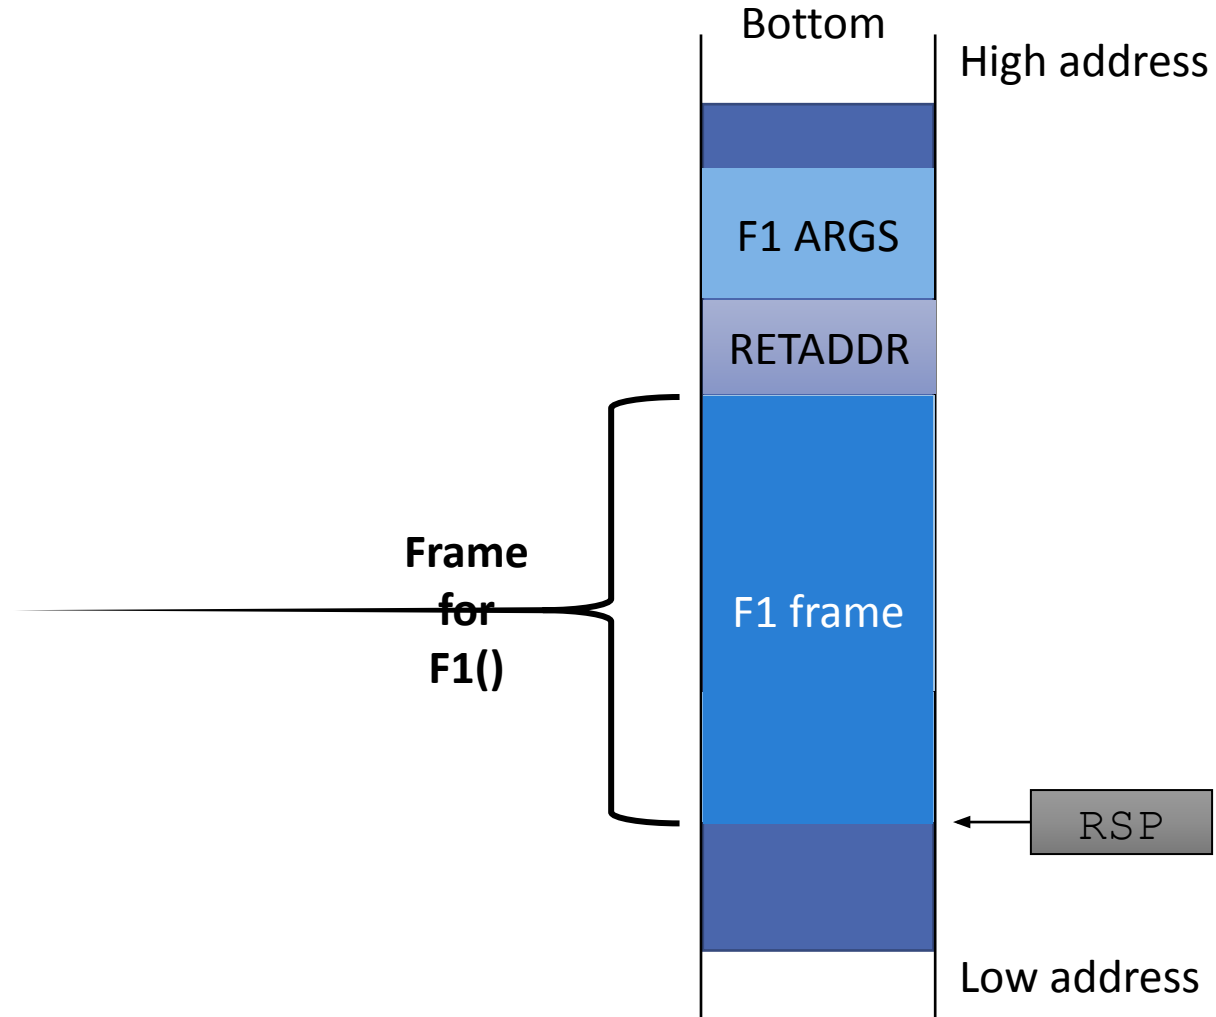
Stack Frame Boundaries

- Start below return address
- Stop at stack pointer

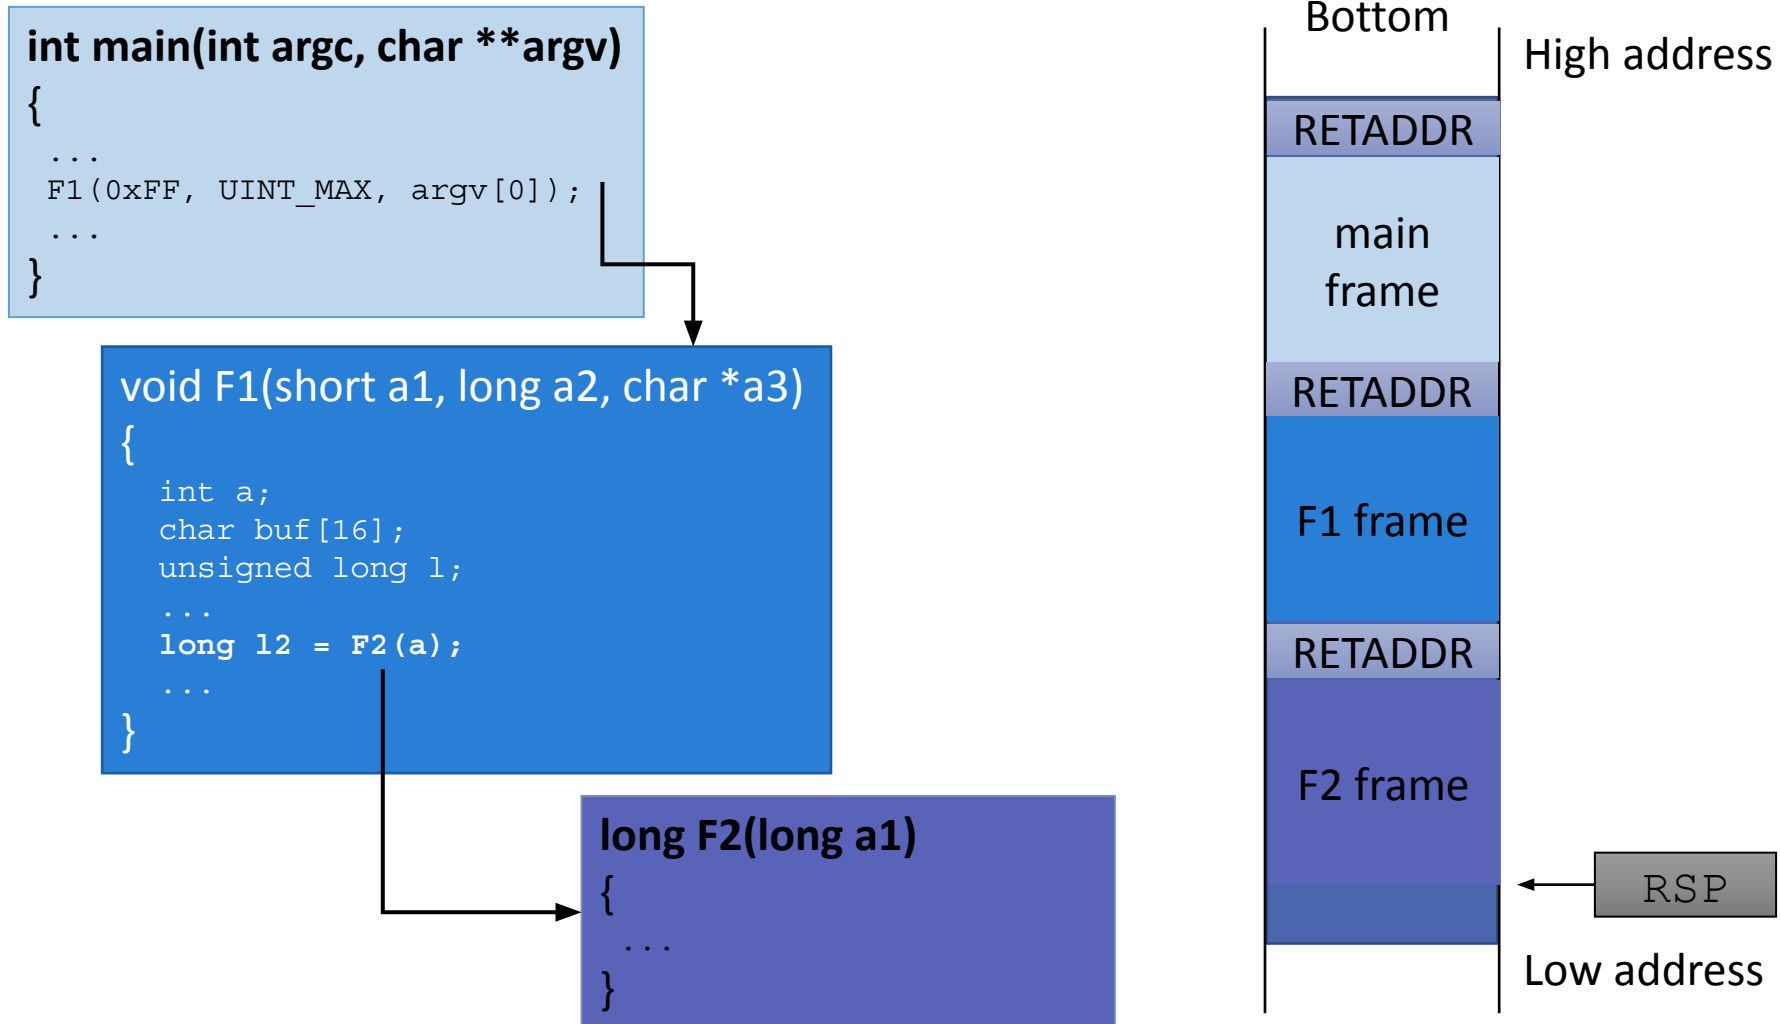


RETADDR and Stack Frames

- The return address may also be considered part of the frame
- We will not for simplicity

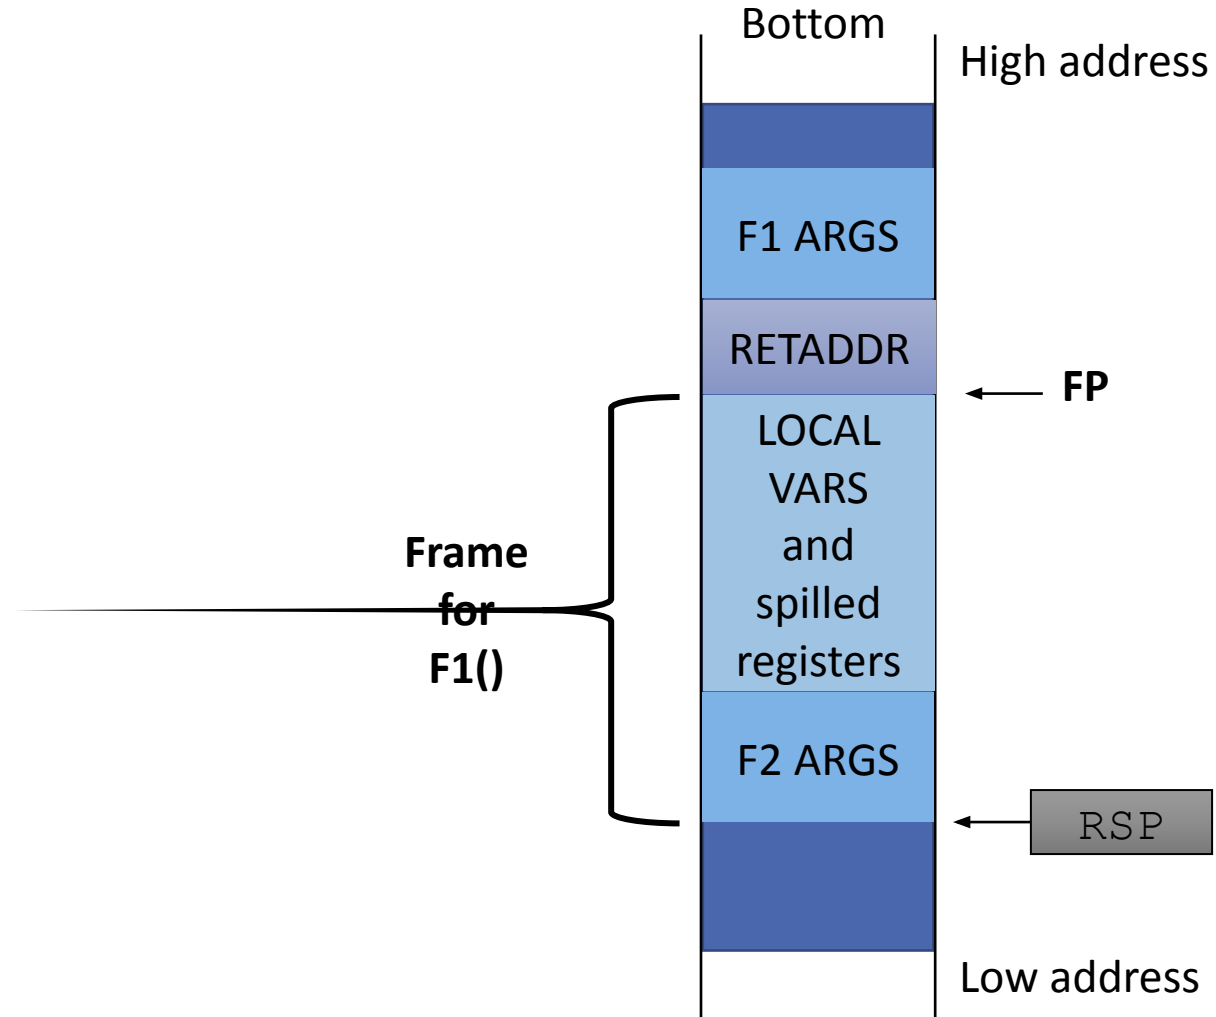


Stack Frames Example



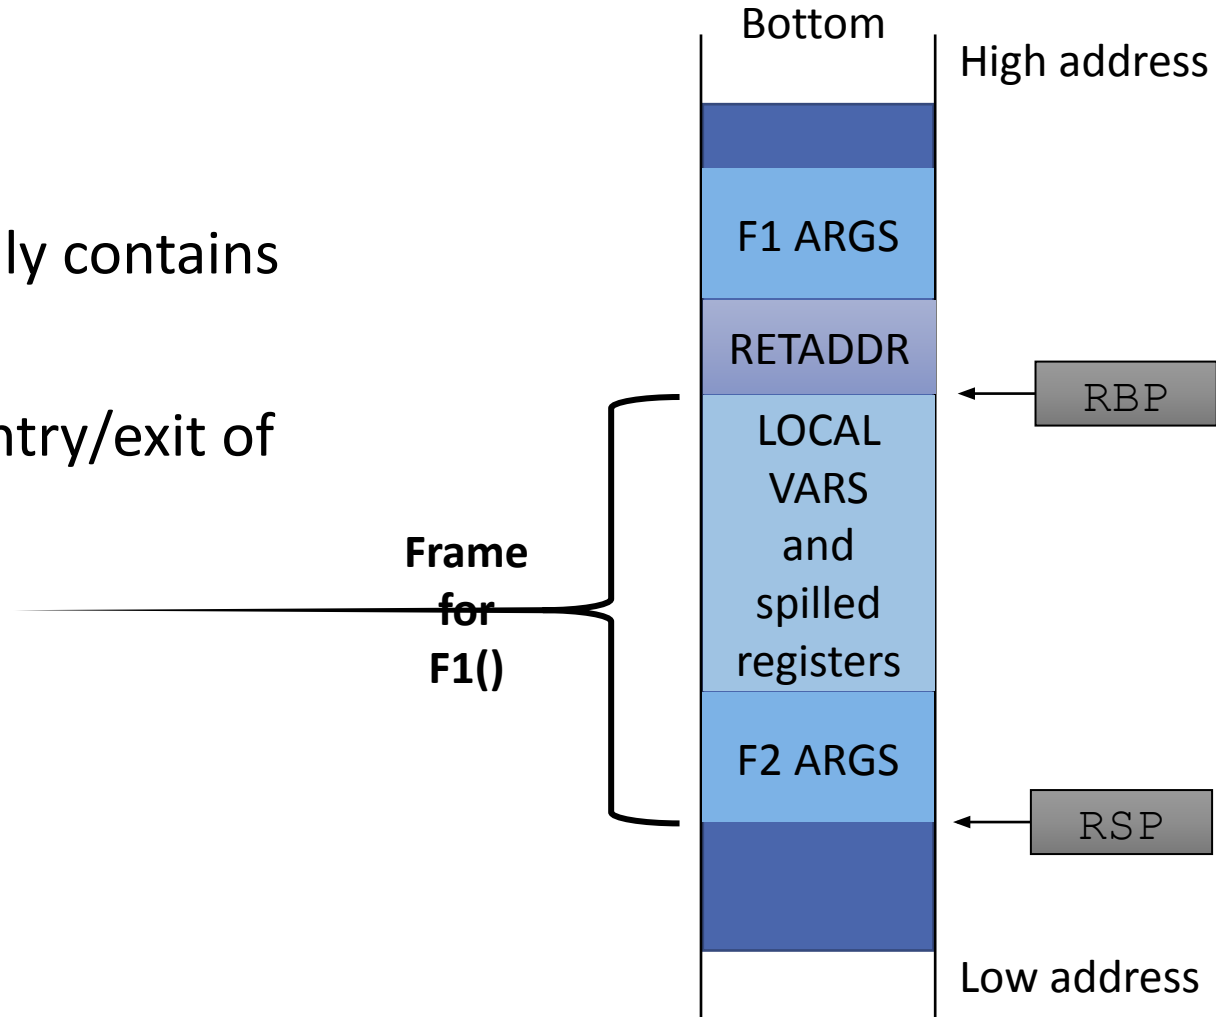
Frame Pointer (FP)

- Marks the highest address in the frame
 - Bottom of the frame
- Aka Base Pointer

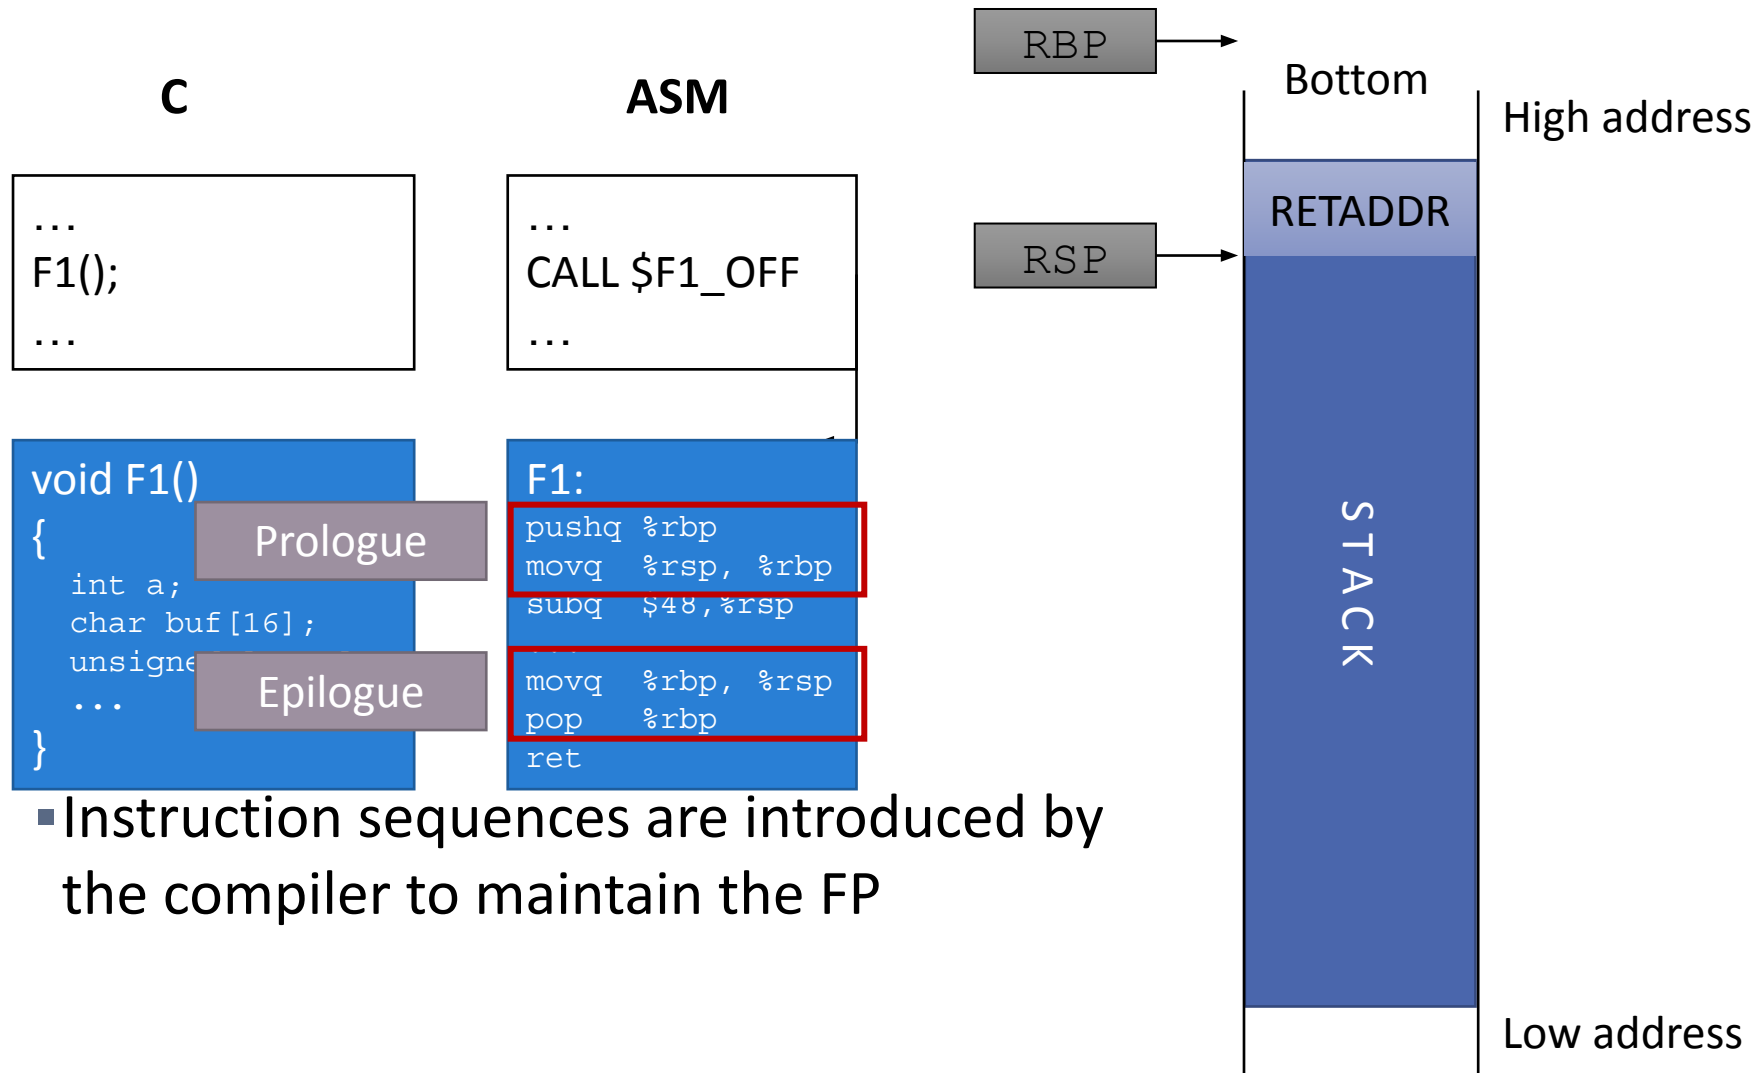


Frame Pointer (FP)

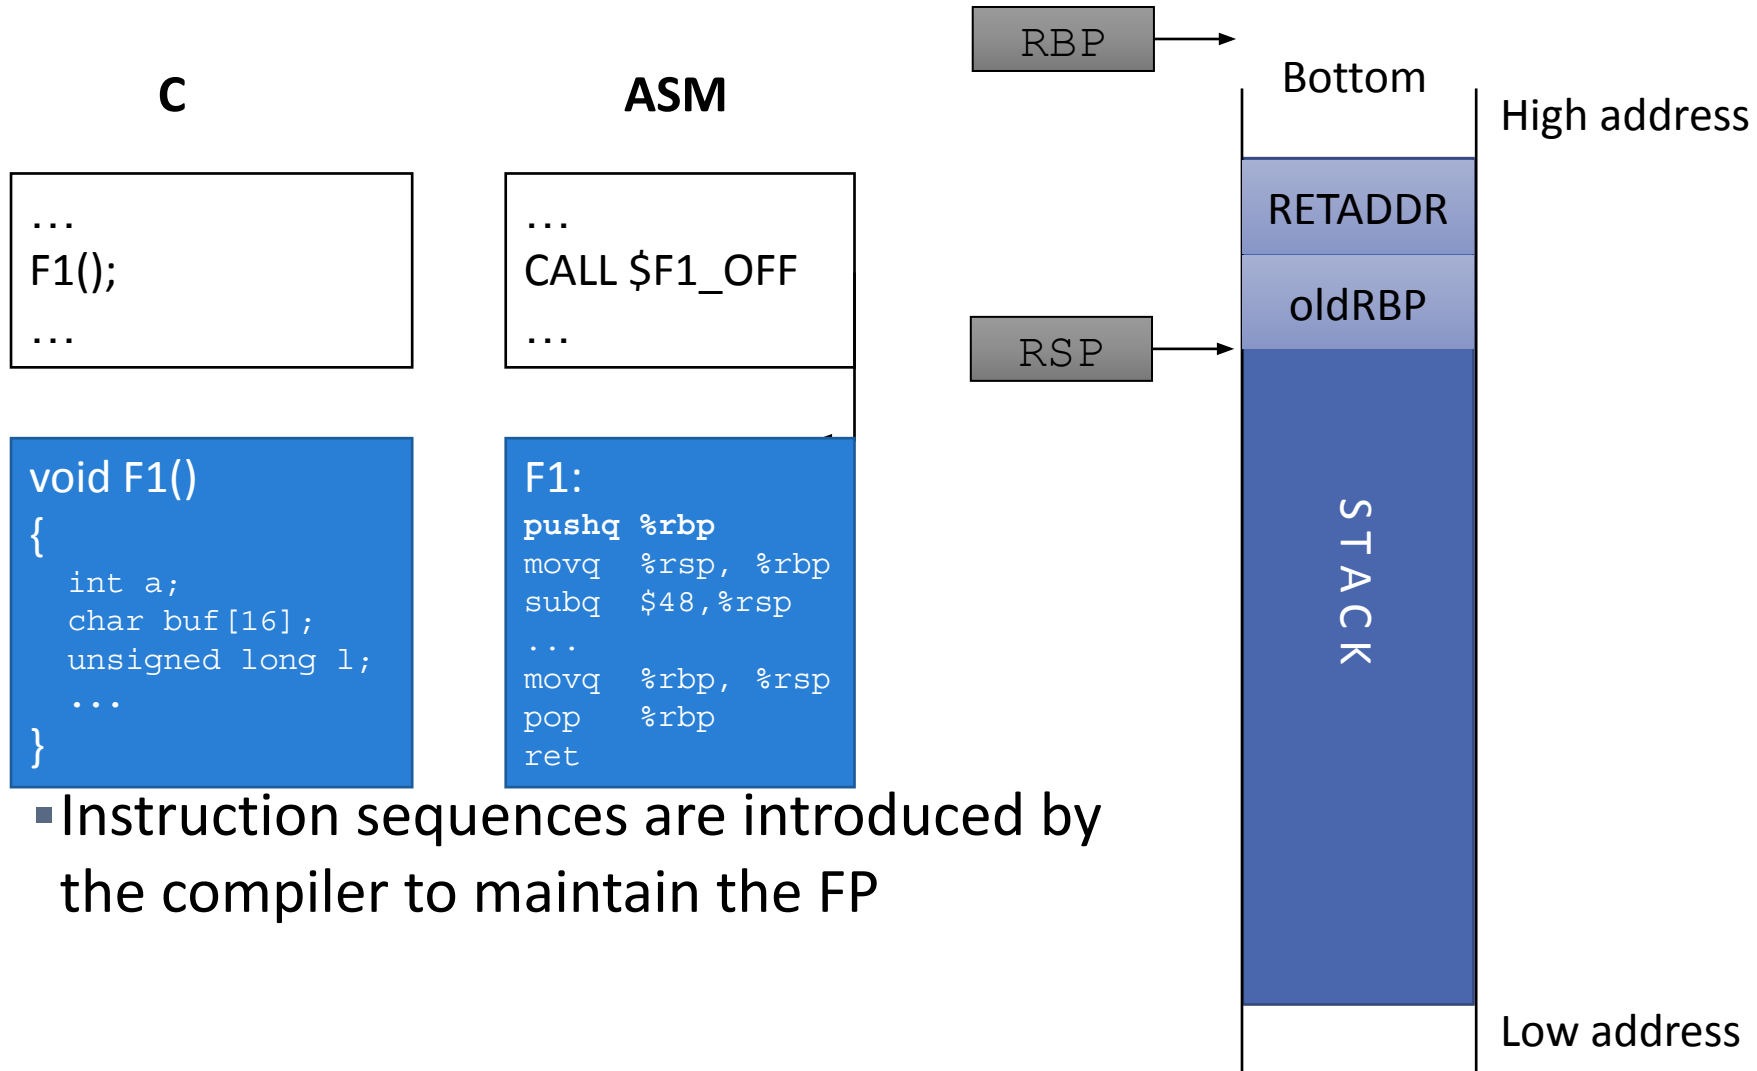
- Marks the highest address in the frame
 - Bottom of the frame
- Aka Base Pointer
- The RBP/EBP register commonly contains the FP
- RBP needs to be updated upon entry/exit of function



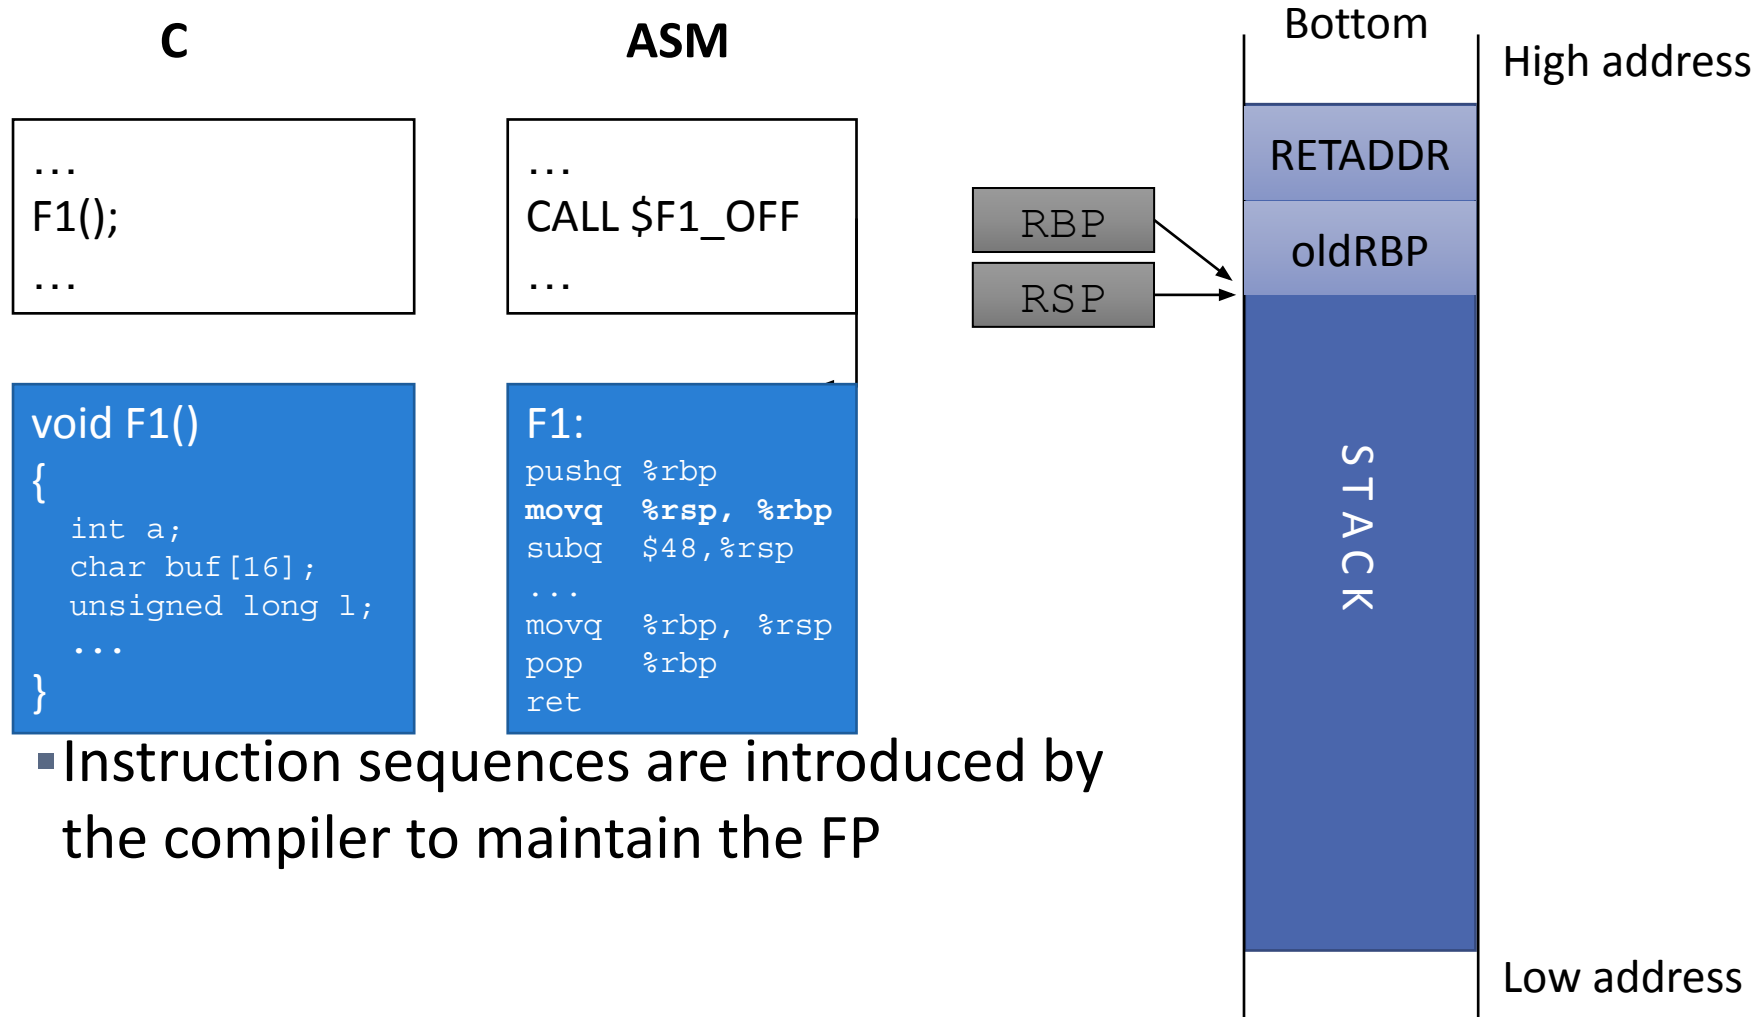
Maintaining the FP



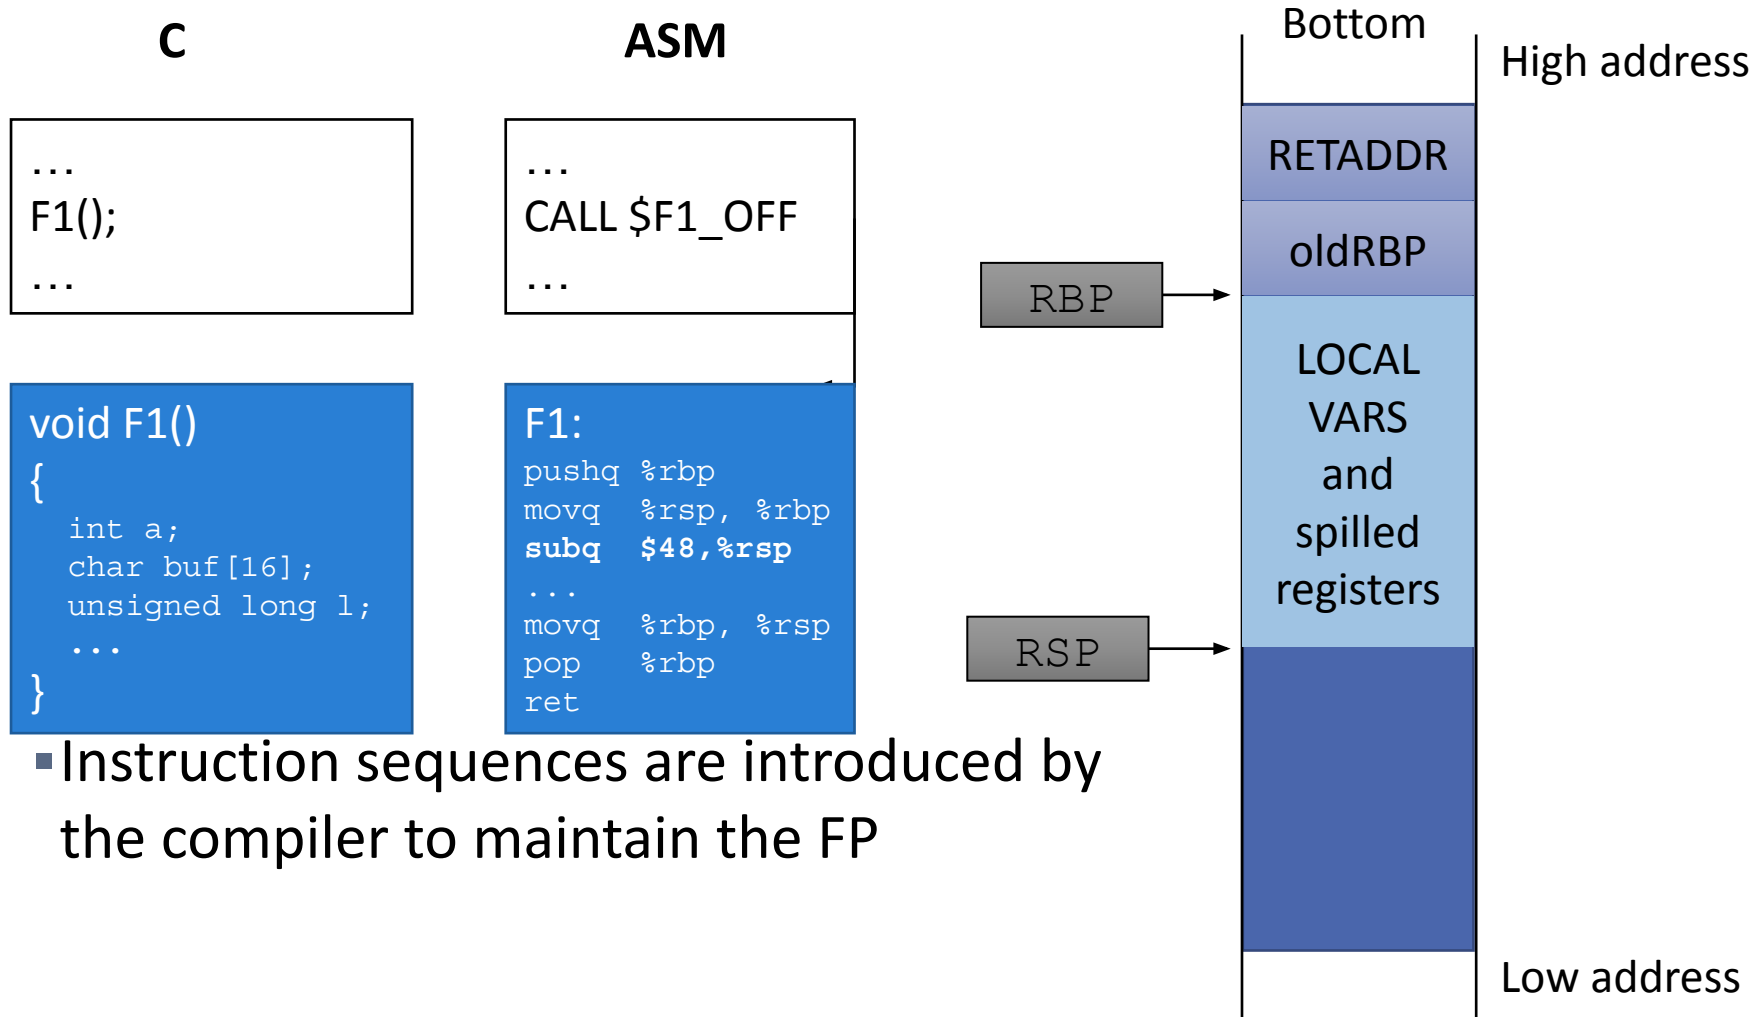
Maintaining the FP



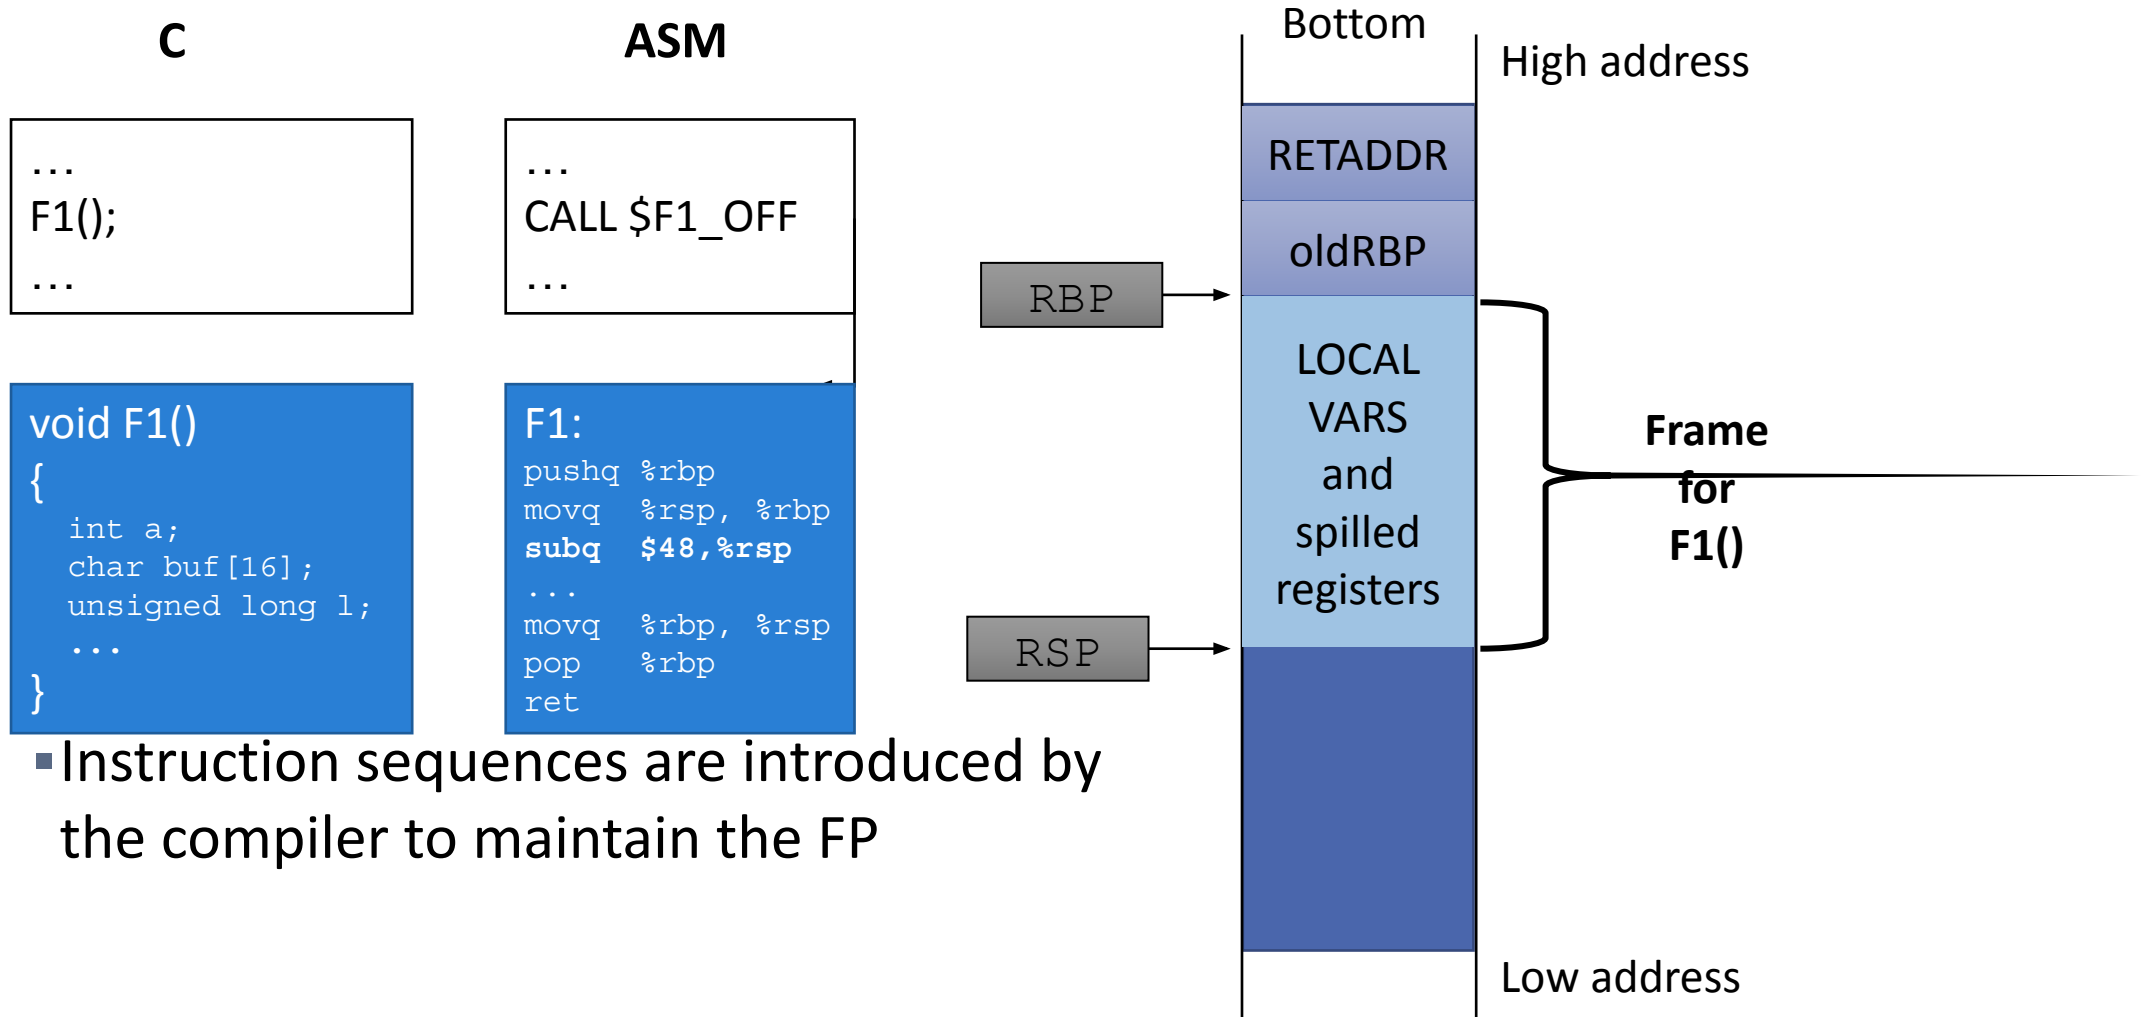
Maintaining the FP



Maintaining the FP

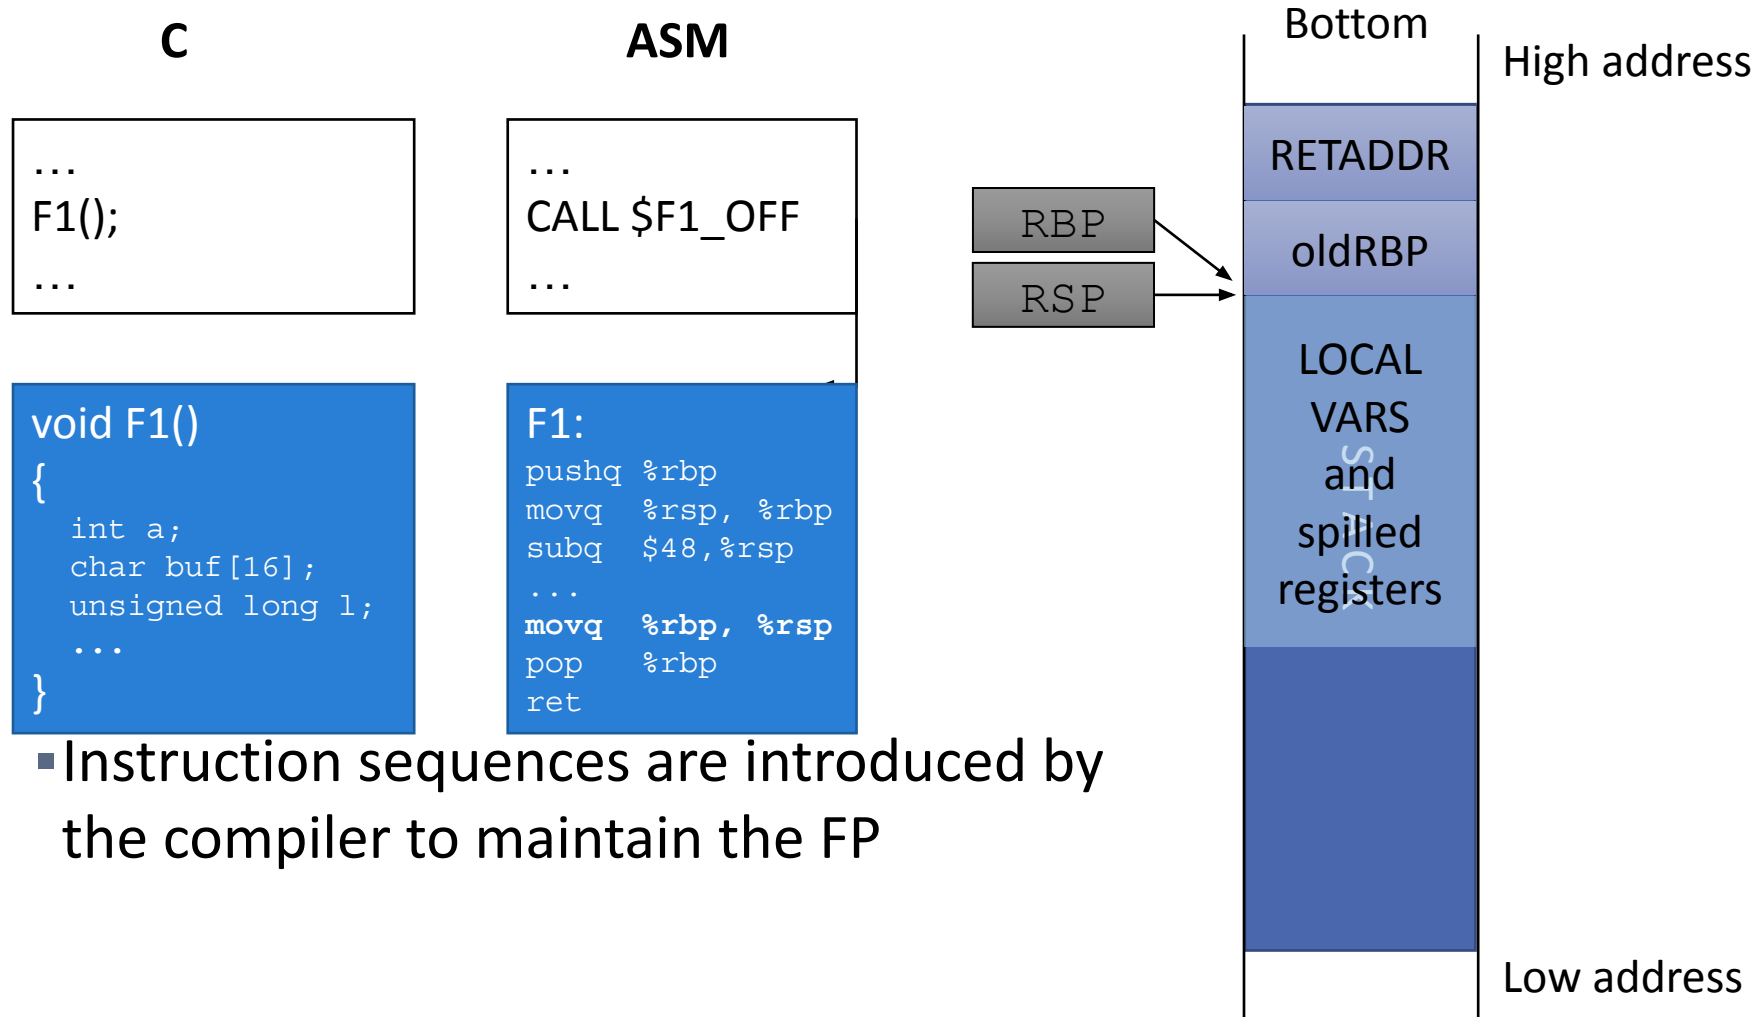


Maintaining the FP

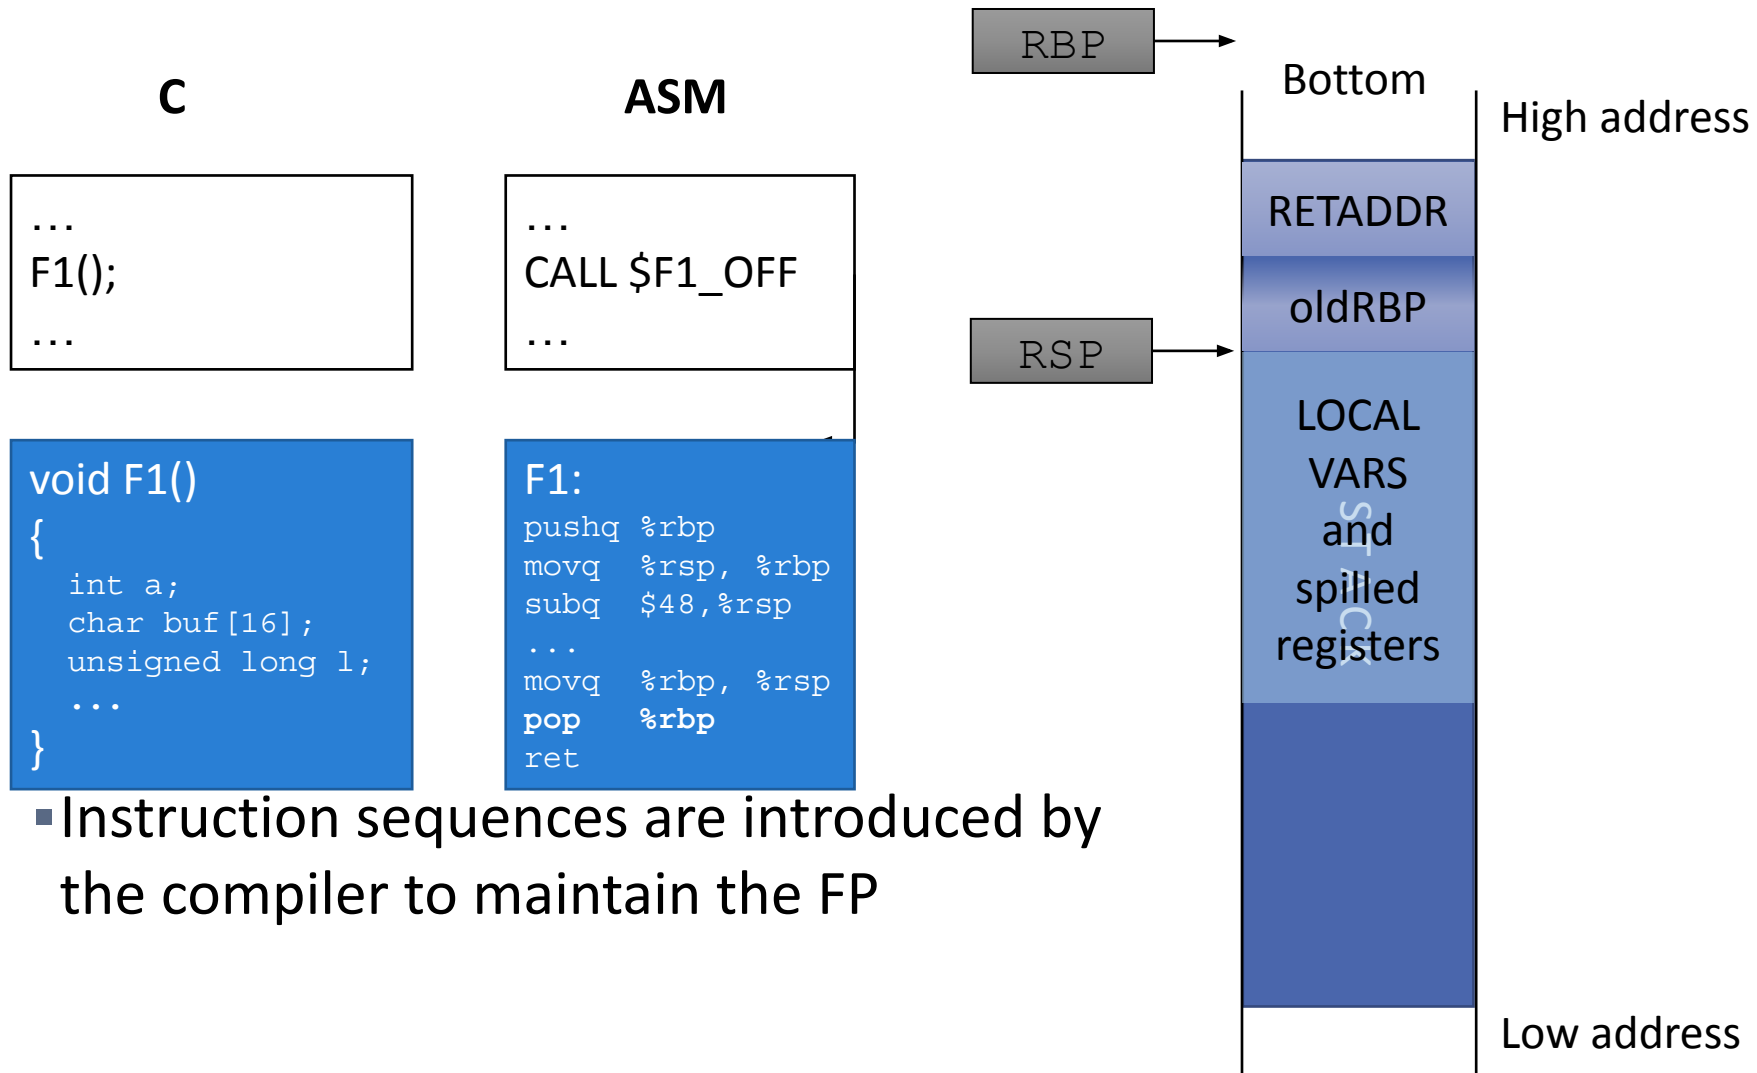


- Instruction sequences are introduced by the compiler to maintain the FP

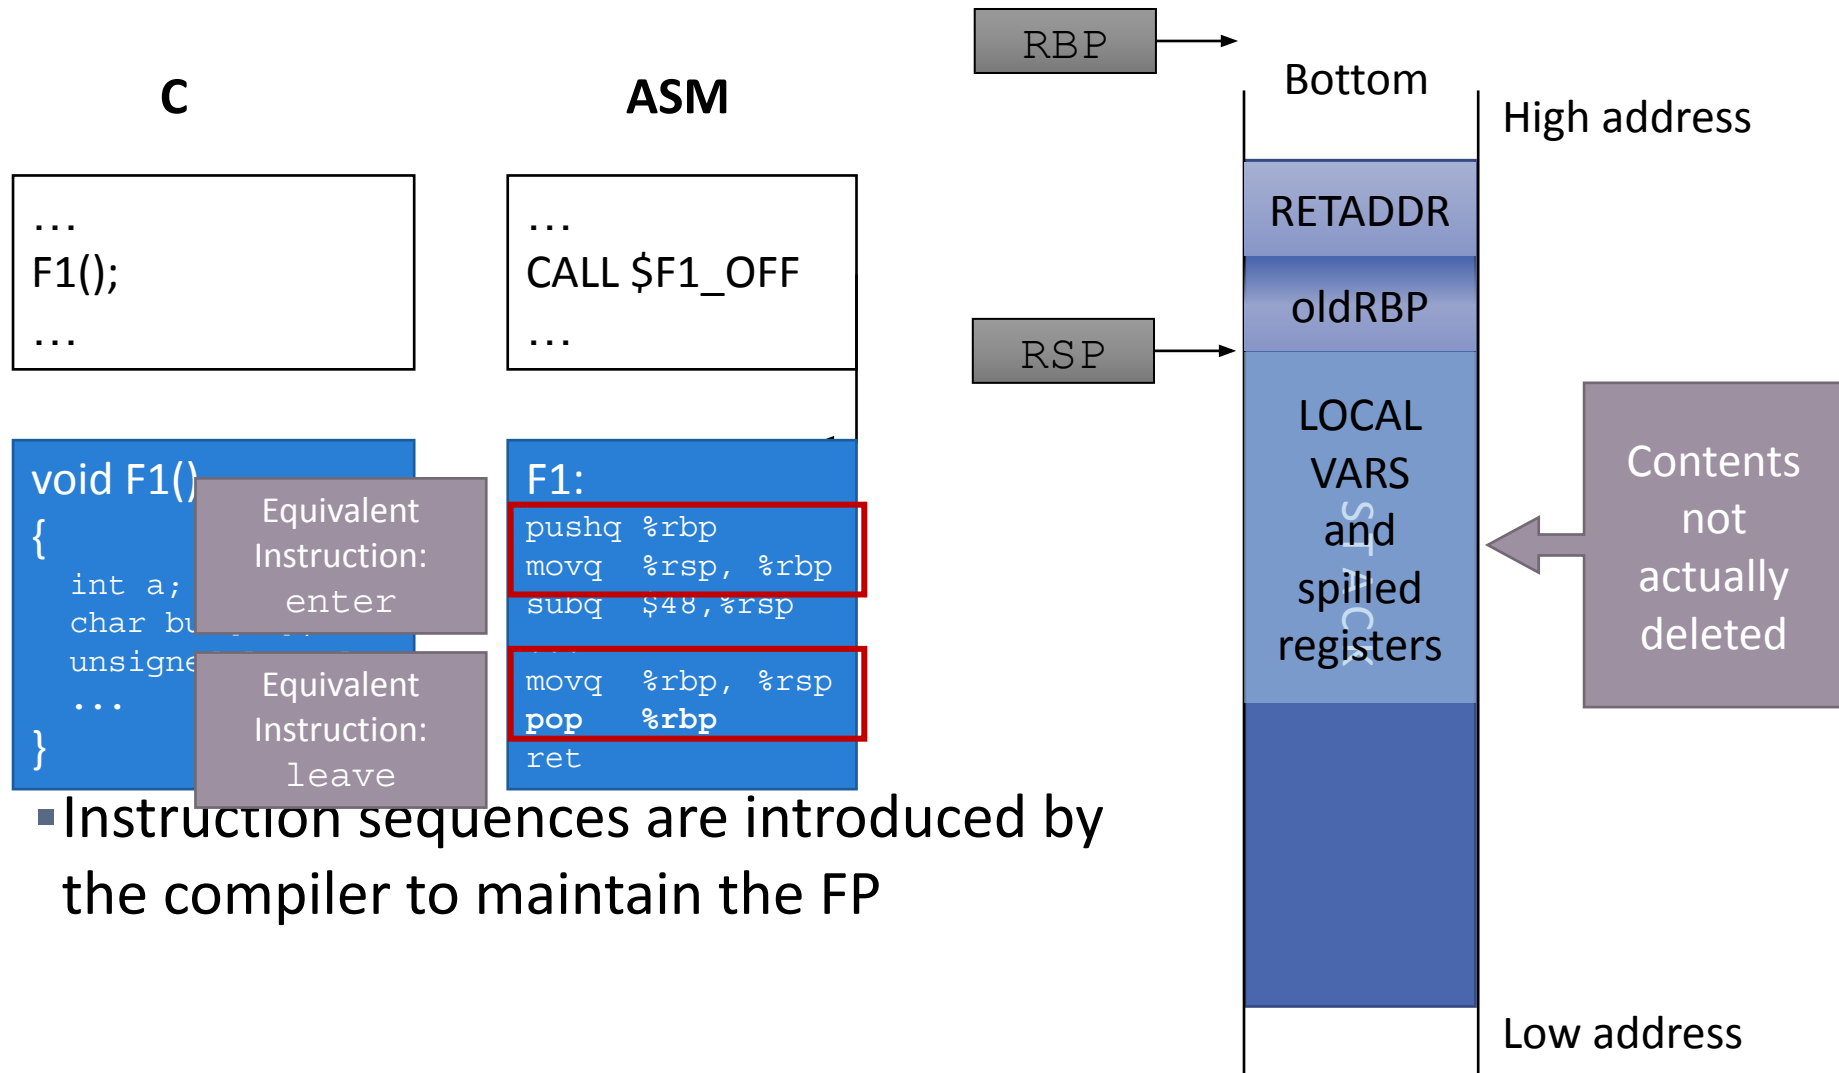
Maintaining the FP



Maintaining the FP



Maintaining the FP



Putting It All Together

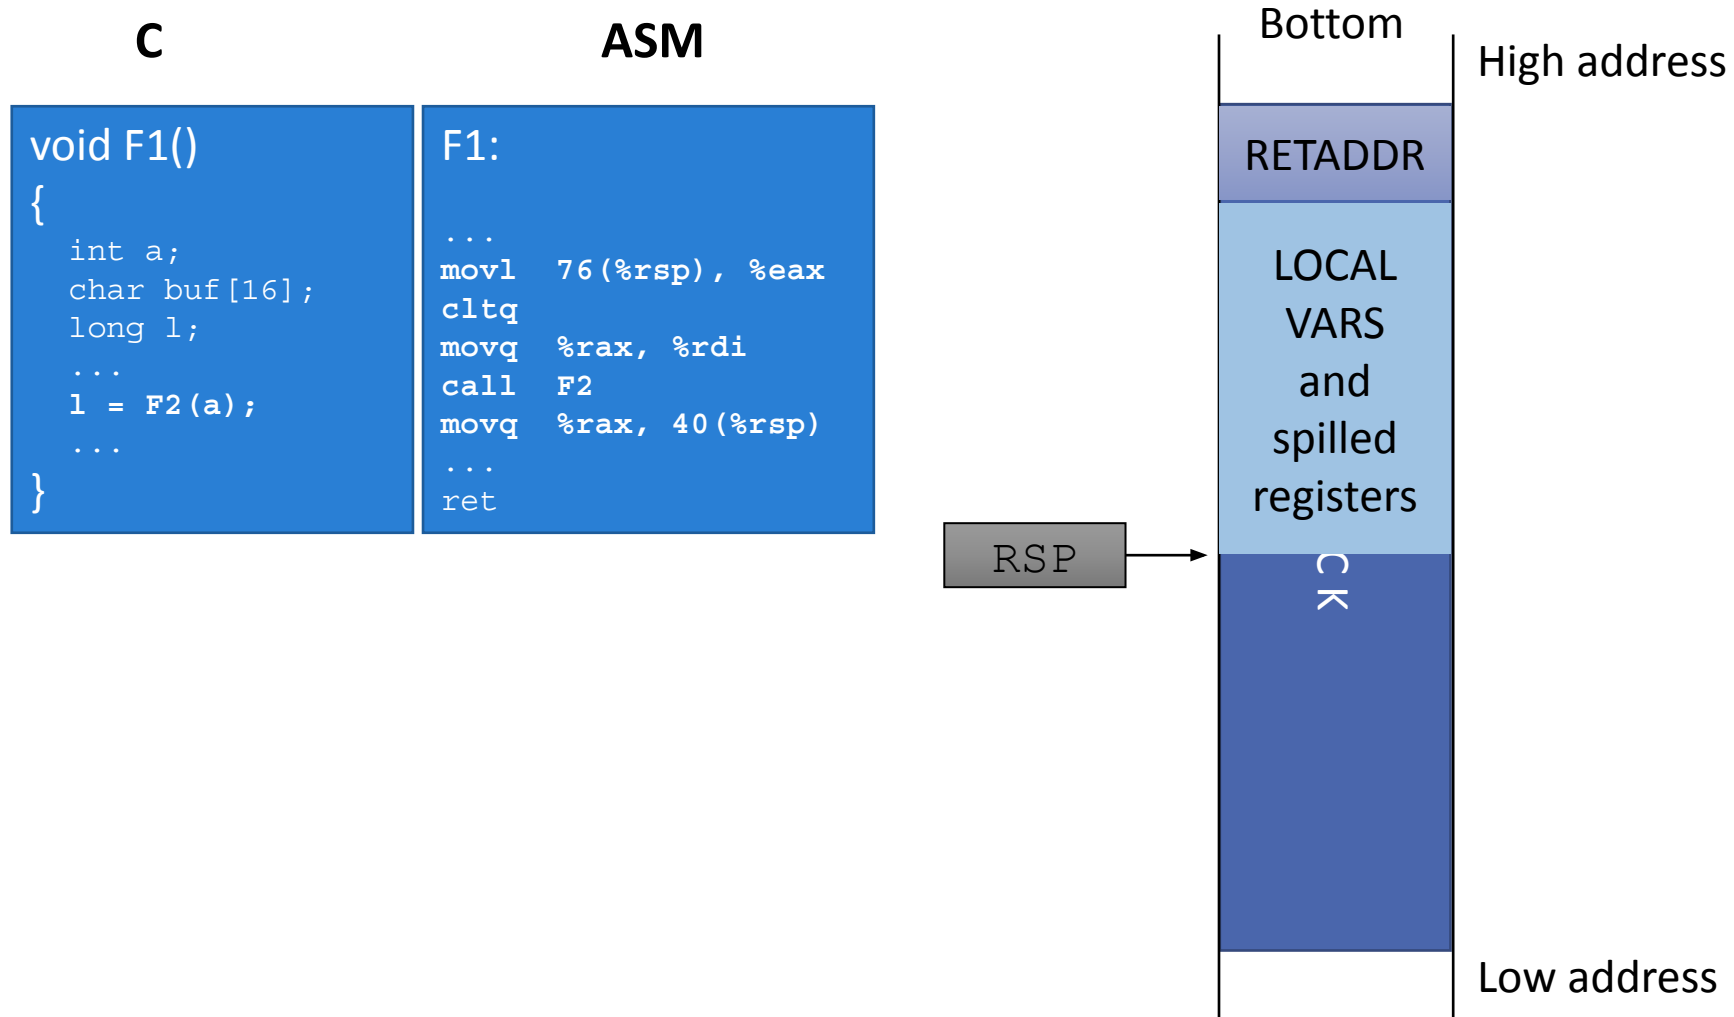
Function Call

- Prepare function call arguments
- Make the call
- Function prologue
 - Save RBP/EBP
 - Setup new RBP/EBP
- Callee saves registers that need to be preserved
- Callee allocates stack space

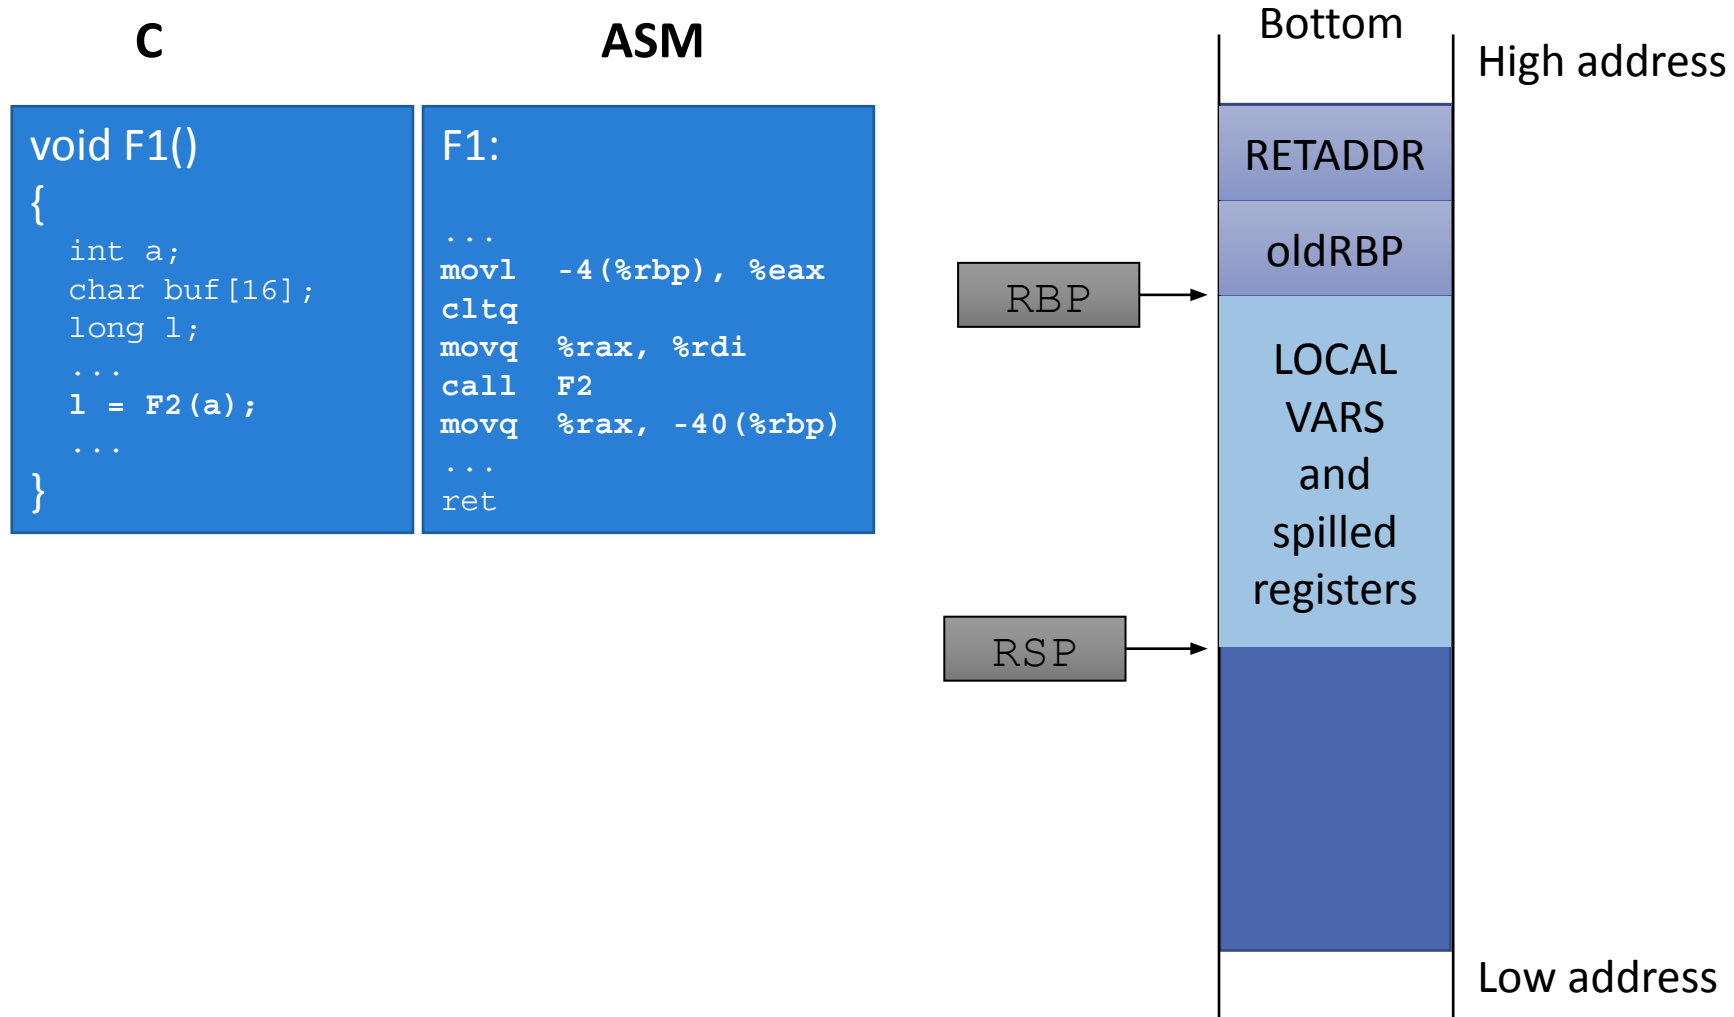
Function Return

- Function epilogue
 - Release stack space
 - Restore BP
- Return

Accessing Stack Variables (no FP)



Accessing Stack Variables (with FP)



Stack Smashing Attacks

Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);


    puts(buf);
}
```

Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```



Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

`./copy AAAAAAAAA`

High address/stack bottom

RETADDR

buf

buf

STACK

Low address/stack top

Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

`./copy AAAAAAAAA`

High address/stack bottom

RETADDR

\0???????

AAAAAAAAA

STACK

Low address/stack top

Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

```
./copy AAAAAAAAAAAAAAAAAAAAAA
```

High address/stack bottom

AAAAAAA\0

AAAAAAA

AAAAAAA

STACK

Low address/stack top

Stack Overflow Example

```
void copy(const char *str)
{
    char buf[16];

    strcpy(buf, str);

    puts(buf);
}
```

```
subq    $40, %rsp
movq    %rdi, 8(%rsp)
movq    8(%rsp), %rdx
leaq    16(%rsp), %rax
movq    %rdx, %rsi
movq    %rax, %rdi
call    strcpy@PLT
leaq    16(%rsp), %rax
movq    %rax, %rdi
call    puts@PLT
nop
addq    $40, %rsp
ret
```

High address/stack bottom

AAAAAAA\0

AAAAAAA

AAAAAAA

STACK

Low address/stack top

Stack Overflow Example

- This stack overflow allows a to control the return address stored in the stack
- When ret executes, the control-flow of the program will be redirected to an arbitrary address □ control-flow hijacking

```
subq    $40, %rsp
movq    %rdi, 8(%rsp)
movq    8(%rsp), %rdx
leaq    16(%rsp), %rax
movq    %rdx, %rsi
movq    %rax, %rdi
call    strcpy@PLT
leaq    16(%rsp), %rax
movq    %rax, %rdi
call    puts@PLT
nop
addq    $40, %rsp
ret
```

High address/stack bottom

AAAAAAA\0

AAAAAAA

AAAAAAA

STACK

Low address/stack top

Control-Flow Hijacking Attacks

- Untrusted inputs that lead to corruption of a code pointer, which will be later dereferenced, lead to **control-flow hijacking attacks**


Original Stack Smashing Attack

- Appeared at Phrack magazine
<http://phrack.org/issues/49/14.html#article>
- Exploits the fact that stack used to be executable
 - Stores binary code in the controlled buffer
 - **Any** executable, controlled buffer will do!
 - Redirect program to inject code
- Performs **arbitrary code injection!**

```
# write(1, message, 13)
mov    $1, %rax          # system call 1 is write
mov    $1, %rdi          # file handle 1 is stdout
mov    $message, %rsi
mov    $13, %rdx         # number of bytes
syscall                     # invoke operating system to do the write


# exit(0)
mov    $60, %rax         # we want return code 0
xor    %rdi, %rdi        # invoke operating system to exit
syscall

message:
.ascii "Hello, world\n"
```



Making Exploits More Robust

- Observation: Different compiler may use different alignment, spill different register, etc.
- Problems:
 - Exact distance of return address may be different between binaries
 - Exact address of buffer may be different



Making Exploits More Robust

- Observation: Different compiler may use different alignment, spill different register, etc.

- Problems:

- Exact distance of return address may be different between binaries
- Exact address of buffer may be different

- Solutions:

- Use multiple copies of the target address
- Prepend a **NOP sled** to shellcode
 - NOPs □ No operations are special one byte instructions to do nothing
- Aim for target address pointing into NOP sled
 - Execution will slide into shellcode

